Classification of bipartite observables I: stable and uniformly stable observables

Eduardo Dueñez and José Iovino

Department of Mathematics, The University of Texas at San Antonio, One UTSA Circle, San Antonio, TX 78249, USA

This is the first of a series of three papers that bring model theory to quantum information. The dividing lines of Shelah's classification theory, which separate the tame first-order theories from the wild ones, become properties of bipartite observables: properties of the expectation values obtained when two independently prepared quantum systems are idealized one after the other. In this paper we focus on stability. An observable is stable when the order in which the two systems are idealized never changes its expectation value, and this is Shelah's stability. Its classical characterizations, in model theory and in analysis, acquire a physical meaning: idealized preparations are simulated by finite mixtures of actual ones, and adaptive probes can tell apart only polynomially many preparations.

Mathematics Subject Classification (2020). Primary 81P15, 81P16, 47B07, 03C45; Secondary 03C20, 03C66, 03C98, 28A35, 46B08, 46B10, 46B50, 46L06, 46L30, 46N50, 68Q11, 68Q32, 81P10, 81P18, 81P45.

Keywords. Order of idealization, bipartite observables, tensor products of states, swap test, Hong–Ou–Mandel interference, simulation by mixtures, independent product, double limit property, weak compactness, Arens regularity, finitely additive measures, Harrington's lemma, model-theoretic stability, decision trees, Littlestone dimension, ultrapowers.

Contents

1 Introduction: one property, many languages **3**

2 Physics: two orders of idealization **12**
2.1 Actual states and effective observables 12
2.2 Idealized states 13
2.3 The two orders of idealization 13

3 Analysis: stability is the double limit property **15**
3.1 Historical remarks 17

The latest revision of this paper is available at https://utsa-quantum-math.github.io/resources/bipartite-observables-1-stability/.

Eduardo Dueñez: Eduardo.Duenez@utsa.edu
José Iovino: jose.iovino@utsa.edu

4 Experiments: the swap test, interference and equality tests are stable **17**
4.1 The minimal tensor product . . . 17
4.2 Factorization through a Hilbert space . . . 18
4.3 The swap operator, symmetrization, and Hong–Ou–Mandel interference . . . 19
4.4 Equality and order of discrete spectra . . . 20
4.5 The order projection . . . 20

5 Topology: idealized responses are observables **22**
5.1 Topology: weak compactness . . . 22
5.2 Geometry: factorization through a reflexive space . . . 24
5.3 Algebra: the two orders of idealization agree . . . 24
5.4 Pure probes suffice . . . 25
5.5 Historical remarks . . . 25

6 Convexity: finite mixtures simulate idealized preparations **26**
6.1 A dimension-free Monte Carlo bound . . . 28
6.2 Historical remarks . . . 29

7 Probability: Fubini's theorem for finitely additive measures **29**
7.1 Fubini's theorem for means . . . 29
7.2 Countably additive measures . . . 30
7.3 Ultrafilter products . . . 30
7.4 Means, mixtures and samples . . . 31
7.5 Historical remarks . . . 31

8 Independence: Harrington's lemma and the independent product **31**
8.1 Abstract independence . . . 31
8.2 Harrington's lemma as an exchange of limits . . . 32
8.3 The independent product of idealized preparations . . . 33
8.4 Historical remarks . . . 34

9 Combinatorics: polynomial versus exponential growth **35**
9.1 From ladders to decision trees . . . 35
9.2 Ultrapowers and uniform stability . . . 37
9.3 Counting the preparations resolved by adaptive probes . . . 40
9.4 Affine approximation versus exponential growth . . . 42
9.5 Learning: Littlestone's mistake bound . . . 43
9.6 Historical remarks . . . 43

10 Order: stable observables are closed under norm limits and orthogonal sums **44**
10.1 Topological and algebraic closure . . . 44
10.2 The boundary of closure: meets and strong limits . . . 45

11 Outlook: registers of growing size **45**

12 Questions **46**

1 Introduction: one property, many languages

Idealized limits, such as infinite energy, vanishing bandwidth or the thermodynamic limit, in which the number of particles and the volume grow without bound at fixed density, are routine in physics, and when two of them are taken in succession their order can matter: in spontaneous symmetry breaking, for instance in a ferromagnet, the thermodynamic limit and the limit of vanishing external magnetic field do not commute. We study the order of idealization for two quantum systems prepared independently and idealized separately, and determine when the expectation of a bipartite observable, that is, of an observable of the joint system, is independent of which system is idealized first. This independence of the order and a uniform form of it govern three operational tasks: the simulation of idealized preparations by finite mixtures, the independent combination of idealized preparations of the two systems, and the discrimination of preparations by adaptive measurements, together with the prediction of the outcomes of these measurements.

Let H_1 and H_2 be complex separable Hilbert spaces describing two quantum systems, which we call Alice's and Bob's. We write $\mathfrak{D}(H_1)$ and $\mathfrak{D}(H_2)$ for the sets of density operators on H_1 and H_2, the positive operators of trace one, which describe the states, or preparations, of the two systems. Measurable quantities of the joint system are represented by self-adjoint operators on $H_1 \otimes H_2$, but we call every bounded operator on $H_1 \otimes H_2$ an observable. When the two systems are prepared independently in states $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$, the joint system is described by $\rho \otimes \sigma$, and an observable C has expectation

$$\varphi_C(\rho, \sigma) = \mathrm{Tr}\big((\rho \otimes \sigma)C\big).$$

Calling every bounded operator an observable costs nothing: the expectation of the adjoint C^* is the complex conjugate of that of C, and the self-adjoint operators $\frac{1}{2}(C + C^*)$ and $\frac{1}{2i}(C - C^*)$ therefore have as expectations the real and imaginary parts of that of C. When the observable is an orthogonal projection, which represents a yes/no test, its expectation is the probability that the test answers yes.

Two orders of idealization. Asymptotic regimes are described by sequences of preparations: states of increasing energy, outputs of circuits of increasing depth, single-photon states of increasingly narrow bandwidth. Such sequences need not converge to a density operator: if Alice's system is an oscillator, for instance, the sequence of its energy eigenstates has no limit among density operators. In the weak* topology, in which states converge when all their expectations converge, the limit points of this sequence are states of the algebra $B(H_1)$ of bounded operators on H_1, that is, positive linear functionals on $B(H_1)$ that take the value one at the identity, and none of them is given by a density operator. We call a state that is not given by a density operator *idealized* (see Section 2); idealized states play the role that plane waves and position eigenstates play in textbook quantum mechanics. Idealizing both systems requires two limits, which can be taken in either order. We call an observable C *stable* if for all sequences $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $(\sigma_l)_{l\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$ we have

$$\lim_{l\to\infty} \lim_{k\to\infty} \varphi_C(\rho_k, \sigma_l) = \lim_{k\to\infty} \lim_{l\to\infty} \varphi_C(\rho_k, \sigma_l) \tag{1}$$

whenever both iterated limits exist. The term comes from model theory and Banach space theory, where versions of the same condition, imposed on formulas and on norms, respectively, define stability [She71, KM81]; it is unrelated to dynamical stability or to stability of matter. An instability occurs already for an observable that is diagonal in a product basis, and hence classical: let $(e_k)_{k\in\mathbb{N}}$ and $(f_l)_{l\in\mathbb{N}}$ be orthonormal sequences in

H_1 and H_2, for instance energy eigenstates of two identical oscillators, and let the *order projection* $P_{k<l}$ be the orthogonal projection onto the closed span of the vectors $e_k \otimes f_l$ with $k < l$. Writing e_k also for the orthogonal projection onto the line through e_k, a density operator, and likewise for f_l, we find that the value $\varphi_{P_{k<l}}(e_k, f_l)$ is 1 if $k < l$ and 0 if $k \geq l$, so, for the sequences $\rho_k = e_k$ and $\sigma_l = f_l$, the left side of (1) is 0 and the right side is 1: the event "Bob's system has the higher energy" has no well-defined probability once both energies are idealized to infinity.

Three operational tasks. For an observable C, a preparation σ of Bob's system acts as a *probe* of Alice's system: it assigns to each preparation ρ of Alice's system the *response* $\varphi_C(\rho, \sigma)$. Dually, once Alice has prepared ρ, the observable C acts on Bob's system as the *effective observable* $\mathrm{Tr}_1((\rho \otimes I)C)$, the operator on H_2 obtained from $(\rho \otimes I)C$ by the partial trace over Alice's system, whose expectation in every preparation σ of Bob's system is the response $\varphi_C(\rho, \sigma)$. The order of idealization bears on three operational tasks:

(i) *Simulation.* An idealized preparation is not available in the laboratory, but actual preparations, those given by density operators, are available, and so are their finite mixtures, in which Alice chooses one of finitely many actual preparations at random with prescribed probabilities. Once an order of idealization is fixed, the task is to reproduce the responses of the idealized preparation, to any prescribed accuracy, by those of a single finite mixture, which must serve for every probe, idealized probes included.

(ii) *Independent combination.* Given an idealized preparation of each system, one seeks an extension of φ_C to such pairs that is weak* continuous in each preparation separately and symmetric in the two systems.

(iii) *Discrimination and learning.* Bob learns about an unknown preparation of Alice's system online, probe by probe: he chooses each probe according to the previous answers and finds out whether the response, estimated by repeating the measurement, lies above or below a fixed threshold. For a positive integer n, a protocol with n probes has 2^n possible answer sequences, so in principle n probes can tell apart exponentially many preparations. What matters is how the largest number of preparations that n probes can tell apart grows with n, exponentially or only polynomially, and how many mistakes Bob makes if he predicts each answer before he receives it.

In Theorems 6.1, 5.1 and 8.2 we show that each of the tasks (i) and (ii) can be carried out if and only if the observable is stable, task (i) with the probing system idealized first. For (ii), stability is equivalent to the agreement of the two orders of idealization on every pair of idealized preparations, one of each system. On the stable observables, the common values define the independent product of the two preparations; it is symmetric in the two systems, and for each stable observable it gives the only extension of the expectation that is weak* continuous in each preparation separately. Task (iii) requires more: as we show in Theorem 9.9, for a self-adjoint observable the growth is polynomial at every margin (the width of a band around the threshold inside which answers are not counted) exactly when the observable satisfies a uniform form of stability, which is strictly stronger than stability (see Theorem 9.6 and Example 9.8).

Recall that, for a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, every bounded sequence $(a_k)_{k\in\mathbb{N}}$ of complex numbers has a unique ultralimit $\lim_{k,\mathcal{U}} a_k$, the complex number a with $\{k : |a_k - a| < \varepsilon\} \in \mathcal{U}$ for every $\varepsilon > 0$; it is the ordinary limit whenever the latter exists. For a

sequence $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ we write $\lim_{k,\mathcal{U}} \rho_k$ for the state $A \mapsto \lim_{k,\mathcal{U}} \operatorname{Tr}(\rho_k A)$ of $B(H_1)$. In Theorem 9.6 we prove that, for a self-adjoint observable C, uniform stability, which we define in Section 9.2, is equivalent to the stability, for every (equivalently, for some) nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, of the ultrapower expectation $\Phi_C^{\mathcal{U}}(x, y) = \lim_{k,\mathcal{U}} \varphi_C(x_k, y_k)$, a function of sequences $x = (x_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $y = (y_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$, which we regard as families of experiments indexed by k. Here a bounded function of two variables is stable when it satisfies the double limit condition (1). Fix a margin $\delta > 0$, and count an answer only when the response lies at least $\delta/2$ away from the threshold. If a self-adjoint observable C is uniformly stable, then there is an integer $d \geq 0$, depending on δ, such that for every positive integer n, protocols with n probes tell apart at most $\binom{n}{0} + \binom{n}{1} + \cdots + \binom{n}{d}$ preparations, a polynomial in n of degree d (see Theorem 9.9). Bob can then also predict online, with at most d mistakes, on which side of a fixed threshold the response of an unknown preparation lies (see Proposition 9.13). If a self-adjoint observable is not uniformly stable, then at some margin, protocols with n probes tell apart 2^n preparations for every positive integer n, and no bound on the number of mistakes holds uniformly in the threshold. Binary decision trees, which describe the adaptive protocols of (iii), connect uniform stability with these counts (see Section 9): uniform stability bounds their height at margin δ by the number d above, a counting argument of the kind that proves the Sauer–Shelah lemma on families of sets turns this bound into the polynomial, and the failure of uniform stability yields, at some margin, trees of every height, that is, for every positive integer n, protocols with n probes that realize all 2^n answer sequences.

The counting has an infinite precedent in Rowbottom's original definition of stability (see Remark 9.10). A complete first-order theory, that is, a maximal consistent set of first-order sentences, is stable in an infinite cardinal λ if, in every model of the theory (every structure in which its sentences hold) and over every set of λ parameters from the model, there are at most λ complete 1-types, that is, complete descriptions of a single element over those parameters. This is the least possible number, whereas a priori there can be 2^λ. Since $\lambda^d = \lambda$ for every integer $d \geq 1$, the distinction between polynomial and linear growth disappears on infinite sets.

Experiments. In Section 4 we show that stability holds for several basic observables of quantum optics and quantum information:

- *Product and local measurements* $A \otimes B$ and $A \otimes I + I \otimes B$, respectively, for bounded operators A on H_1 and B on H_2; operator-norm limits of finite sums of such products (the minimal tensor product); all compact (for instance, finite-rank) observables, such as the projections onto Bell states, maximally entangled states of two qubits (quantum two-level systems); and all observables when one of the systems is finite-dimensional (see Theorem 4.1 and Corollary 4.2).

- *The swap test* [BCWdW01], for two systems with the same Hilbert space: its acceptance probability $\frac{1}{2}(1 + \operatorname{Tr} \rho\sigma)$ determines the overlap $\operatorname{Tr} \rho\sigma$ of two unknown states, and its observable, the projection onto the symmetric subspace, is stable, as are the projection onto the antisymmetric subspace and the swap operator, which exchanges the two systems (see Section 4.3).

- *Hong–Ou–Mandel coincidences* [HOM87]: a balanced (50:50) beam splitter, such as a half-silvered mirror, transmits each single incoming photon with probability $\frac{1}{2}$ and reflects it otherwise; when two single photons whose internal (spectral and polarization) states are ρ and σ enter its two ports, a coincidence, a detection at both output ports, occurs with probability $\frac{1}{2}(1 - \operatorname{Tr} \rho\sigma)$ (see Section 4.3).

- *Discrete equality tests*, the projections onto the event that two quantities with pure point spectra, that is, with an orthonormal basis of eigenvectors, take equal values, such as "equal photon numbers" in two modes of light or "equal energies" of two particles confined in identical boxes (see Proposition 4.6).

We also show that, for infinite-dimensional systems, the swap operator and the projections of the swap test and of Hong–Ou–Mandel coincidences are not norm limits of finite sums of products, and neither are the equality tests between quantities with infinitely many common values (see Proposition 4.3, Section 4.3 and Proposition 4.6). The swap operator, the projections of the swap test and of Hong–Ou–Mandel coincidences, and the equality tests are nevertheless stable for a simple reason: the expectation of each has the form $\varphi_C(\rho,\sigma) = \langle u(\rho), v(\sigma)\rangle$ for bounded maps u from $\mathfrak{D}(H_1)$ and v from $\mathfrak{D}(H_2)$ into a common Hilbert space; we call such observables *Hilbert-factorable*, and in Proposition 4.4 we show that every Hilbert-factorable observable is stable. For the swap operator, for instance, the expectation $\operatorname{Tr}\rho\sigma$ is the Hilbert–Schmidt inner product of the two density operators. With Hilbert space replaced by a reflexive Banach space, one that the canonical embedding identifies with its second dual, this form characterizes stability: an observable is stable exactly when its expectation has the form $\langle u(\rho), v(\sigma)\rangle$ for bounded affine maps u from $\mathfrak{D}(H_1)$ into a reflexive Banach space E and v from $\mathfrak{D}(H_2)$ into its dual E^* (see Theorem A below, whose condition (A3) is the operator form of a theorem of Iovino [Iov99b]).

By contrast, the comparison $[A < B]$, the projection onto the event that Bob's quantity B takes a larger value than Alice's quantity A, is unstable whenever A and B are self-adjoint operators on H_1 and H_2 with pure point spectra that are unbounded above, as we prove in Proposition 4.6; examples are "Bob's system has more photons" and "Bob's system has the higher energy", the event of the order projection when the vectors e_k and f_l are the energy eigenstates of two identical oscillators. This instability, like that of the order projection, has nothing to do with noncommutativity: both projections are diagonal in a product basis.

Summary of main results.

(1) **Simulation by finite mixtures.** An observable is stable exactly when every idealized preparation of Alice's system can be reproduced to arbitrary accuracy by finite mixtures of actual preparations, uniformly over all probes, idealized ones included, with Bob's system idealized first; with Alice's system idealized first, simulation can succeed even for unstable observables (see Theorem 6.1 and Example 6.2). For the swap test, Hong–Ou–Mandel coincidences and equality tests, every ultralimit of a sequence of actual preparations is reproduced within ε, uniformly over all probes, by an equal-weight mixture of $\lceil 4/\varepsilon^2 \rceil$ terms of the sequence, for every $\varepsilon > 0$ and in every dimension (see Theorem 6.5 and Corollary 6.6).

(2) **Randomization.** An observable is stable exactly when, for any two finitely additive probability measures, one on the actual preparations of each system, the two iterated averages of the expectation agree (a Fubini theorem; see Theorem 7.1); such measures describe random choices whose average states may be idealized, and ordinary (countably additive) randomization of the preparations of either system never detects instability (see Proposition 7.2).

(3) **Independent combination.** An observable is stable exactly when the two orders of idealization assign it the same expectation on every pair of idealized preparations, one of each system; the common values, taken over all stable observables, define

the independent product of the two preparations, a state on the space of stable observables that extends the product of density operators, is symmetric in the two systems and gives, for each stable observable, the only extension of its expectation that is weak* continuous in each preparation separately (see Theorem 8.2).

(4) **Discrimination and learning.** A self-adjoint observable is *uniformly stable*, that is, its ladders, finite staircases of responses like that of the order projection, have bounded length at every margin, exactly when its decision trees have bounded height at every margin. Uniform stability is also equivalent to each of three other conditions: at every margin, the number of preparations that adaptive probes can tell apart grows only polynomially with the number of probes; at every margin, the number of mistakes in predicting online on which side of a fixed threshold the response of an unknown preparation lies has a bound independent of the threshold; and the expectation remains stable when preparations and probes are replaced by families of experiments indexed by $k \in \mathbb{N}$, with the expectation taken as the ultralimit over k along any nonprincipal ultrafilter on $\mathbb{N}$. If a self-adjoint observable is not uniformly stable, then at some margin it encodes decision trees of every height, and for every positive integer n, protocols with n adaptive probes then tell apart 2^n preparations. Uniform stability is strictly stronger than stability: a comparison of levels within spin multiplets, the irreducible representations of the group SU(2) through which rotations act in quantum mechanics, is stable but not uniformly stable (see Theorem 9.6 and Example 9.8). For the swap test and Hong–Ou–Mandel coincidences, trees at a margin $\delta > 0$ have height at most $4/\delta^2$, and for equality tests at most $16/\delta^2$, by Novikoff's perceptron bound (see Section 9).

(5) **Algebraic and topological closure.** Stable observables are closed under linear combinations, adjoints and norm limits, and include the minimal tensor product and all Hilbert-factorable observables. Stable yes/no tests are closed under negation and under combining mutually exclusive tests. Stability is not preserved by the quantum-logical "and", the projection onto the intersection of two ranges, even for two uniformly stable projections, nor by limits in the strong operator topology, that is, by pointwise limits on vectors (see Section 10).

One property, many languages. The double limit condition (1), for bounded functions on a product of two sets, has been found, often independently, in several parts of mathematics: as Grothendieck's criterion for weak compactness [Gro52], as Pták's combinatorial lemma on convex means [Ptá63], as a Fubini property for finitely additive measures [Ptá63, Sim68], as Arens regularity of bilinear forms [Are51, Pym65, Ü87], as Krivine–Maurey stability of Banach spaces [KM81], as Shelah's stability [She71], and as Harrington's lemma on the symmetry of independence [HH84]; its uniform form appears, through Shelah's trees, as finite Littlestone dimension, a bound on the height of decision trees, in learning theory [Lit88, CF19]. Iovino observed that, for two-valued functions, the failure of the double limit condition is Shelah's order property, the existence of an infinite staircase of values like that of the order projection, and, building on Krivine and Maurey, he developed stability for real-valued structures [Iov94, Iov99a, Iov99b].

Model theory as a common language. In real-valued model theory [Kei23], the expectation φ_C of a self-adjoint observable C is the interpretation of a formula $\varphi(x, y)$, a

TWO DICHOTOMIES

Physical question *language; model theory* (section)	**Tame side: stable**	**Wild side: unstable**
Experiments (§4)	swap test, Hong–Ou–Mandel coincidences, equality tests, product and compact observables	energy comparison $[A < B]$ with spectra unbounded above, order projection $P_{k<l}$
DICHOTOMY I: STABILITY (see Theorem A)		
Order of idealization *analysis; order property* (§3)	the two iterated limits agree	infinite ladders: the iterated limits are 1 and 0
Idealized responses *topology; definable types* (§5.1)	effective observables relatively weakly compact; separately continuous extension	an idealized response that is not an observable
Factorization *geometry; reflexive pairing* (§5.2)	expectation $\langle u(\rho), v(\sigma)\rangle$ with bounded affine maps u into a reflexive space E and v into E^*	no factorization, by bounded maps, through a reflexive space
Idealized states *algebra; coheirs in two orders* (§5.3)	Arens regular: the two orders of idealization agree on every pair of states	the two orders of idealization disagree
Simulation *convexity; definitions by averages* (§6)	idealized preparations simulated by finite mixtures, uniformly over all probes, idealized ones included, with the probing system idealized first	simulation error at least 1 for the order projection
Randomization *probability; Keisler measures* (§7)	Fubini for all finitely additive probability measures	Fubini fails for some finitely additive measures, never for countably additive ones
Independent combination *independence; Harrington's lemma* (§8)	a unique, symmetric independent product	the two orders of idealization disagree
DICHOTOMY II: UNIFORM STABILITY (see Theorem B)		
Uniform order of idealization *analysis; order property at a margin* (§9.1)	ladders of bounded length at every margin	ladders of every length at some margin
Decision trees *combinatorics; Shelah's trees* (§9.1)	trees of bounded height at every margin	trees of every height at some margin
Families of experiments *ultraproducts; stability in all models* (§9.2)	expectation on ultrapowers is stable	infinite decision trees in ultrapowers
Discrimination *growth; Rowbottom's λ-stability* (§9.3)	at every margin, the number of preparations that n probes tell apart is polynomial in n	at some margin, protocols with n probes tell apart 2^n preparations for every n
Learning *learnability; Littlestone's mistake bound* (§9.5)	finite mistake bound in the threshold game, uniform in the threshold; at most $4/\delta^2$ at margin δ for the swap test (perceptron)	at some margin, no mistake bound uniform in the threshold

Table 1: Model theory as a Rosetta stone. After the first row, which lists examples, each row states a property of an observable as a physical question, with its mathematical language, its model-theoretic reading (in italics) and the section in which the property is treated; these rows follow the order of the sections. For a bounded observable, either all the tame properties of Dichotomy I hold or all fail (see Theorem A), and the same holds for the properties of Dichotomy II and a self-adjoint observable (see Theorem B). The tame side of Dichotomy II is strictly smaller than that of Dichotomy I (see Theorem 9.6 and Example 9.8). On the wild side, the iterated limits 1 and 0 and the simulation error are those of the sharpest example, the order projection; the other entries hold for every observable on the wild side.

real-valued predicate whose variables x and y range over the preparations of Alice's and of Bob's system, respectively. Each of the physical notions of this Introduction has a standard model-theoretic counterpart, which we describe in Section 8 and in the historical remarks at the end of the sections; the local stability theory on which we draw was developed by Ben Yaacov and Usvyatsov [BYU10] in the continuous logic of Ben Yaacov, Berenstein, Henson and Usvyatsov [BYBHU08]. We take as the model the structure formed by the preparations of the two systems with the predicate φ. An idealized preparation of Alice's system determines a φ-type over the model, that is, an assignment of a value to $\varphi(x, \sigma)$ for every preparation σ of Bob's system. This type is realized in an ultrapower of the model, in the sense that its values are the values $\varphi(a, \sigma)$ at some element a of the ultrapower; for an actual preparation the type is realized in the model itself, and it is given by the effective observable that the preparation leaves on Bob's system. A subset of a Banach space is *relatively weakly compact* if its closure is compact in the weak topology, the topology in which a net converges to a point exactly when its values under every bounded linear functional converge to the value at that point. Relative weak compactness of the set of effective observables corresponds to two properties of types together: their definability, that is, the fact that the values of each type, as a function of the parameters, are given up to uniform approximation by formulas with parameters from the model, and the fact that every extension of a type to idealized parameters that is approximately finitely satisfiable in the model still takes the values given by the definition of the type. Here an extension of a type is *approximately finitely satisfiable* in the model if, for any finitely many of its parameters and every $\varepsilon > 0$, some element of the model has, at each of these parameters, a value of φ within ε of the value that the extension assigns. Definability alone does not suffice, since for the order projection every type determined by an idealized preparation of Alice's system is definable. The two orders of idealization correspond to the two orders of realizing coheirs, the extensions of types to larger sets that are approximately finitely satisfiable in the model, and the agreement of the two orders corresponds to Harrington's lemma, by which the definitions of two types commute. Finitely additive probability measures on the sets of preparations correspond to Keisler measures, the finitely additive probability measures on definable sets; ladders and decision trees correspond to Shelah's order property and to his binary trees of formulas, respectively; and families of experiments correspond to elements of ultrapowers. Under this correspondence, the double limit condition (1) becomes the stability of the formula φ in the model, the absence of infinite ladders there, which is strictly weaker than its stability in the theory of the model; the latter corresponds to uniform stability (see Theorem 9.6 and Example 9.8). In this sense model theory serves as a Rosetta stone: in Table 1, the model-theoretic reading of each row, printed in italics after the semicolon in the first column, is the one through which each row of a dichotomy can be translated into any other row of the same dichotomy. Theorems A and B below state precisely the equivalences among the rows of Dichotomy I and among those of Dichotomy II, respectively.

Theorem A (stability in seven languages). Let H_1 and H_2 be the Hilbert spaces of Alice's and Bob's systems, and let C be a bounded observable on $H_1 \otimes H_2$, with expectation $\varphi_C(\rho, \sigma) = \operatorname{Tr}((\rho \otimes \sigma)C)$ for $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$. Then the following are equivalent:

(A1) *Analysis: order of idealization* (Krivine–Maurey, Shelah). The observable C is stable: for all sequences of preparations of the two systems, the two iterated limits in (1) agree whenever both exist. Definition 3.1

(A2) *Topology: idealized responses* (Grothendieck). The effective observables $\mathrm{Tr}_1((\rho \otimes I)C)$, for $\rho \in \mathfrak{D}(H_1)$, form a relatively weakly compact subset of the Banach space $B(H_2)$; equivalently, the expectation extends to a function on pairs of states, one of each system, that is weak* continuous in each state separately. Theorem 5.1

(A3) *Geometry: factorization* (Iovino; Davis–Figiel–Johnson–Pełczyński). There are a reflexive Banach space E and bounded affine maps u from $\mathfrak{D}(H_1)$ into E and v from $\mathfrak{D}(H_2)$ into the dual E^* such that $\varphi_C(\rho, \sigma) = \langle u(\rho), v(\sigma) \rangle$ for all $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$. Theorem 5.1

(A4) *Algebra: idealized states* (Arens, Pym, Ülger). The two orders of idealization assign the same expectation to C on every pair of idealized states, one of each system. Theorem 5.1

(A5) *Convexity: simulation* (Mazur). Every idealized preparation of Alice's system can be simulated to arbitrary accuracy by finite mixtures of actual preparations, uniformly over all probes, idealized ones included, with Bob's system idealized first. Theorem 6.1

(A6) *Probability: randomization* (Simons, Pták). For every finitely additive probability measure on the subsets of $\mathfrak{D}(H_1)$ and every finitely additive probability measure on the subsets of $\mathfrak{D}(H_2)$, the two iterated integrals of the expectation against them agree. Theorem 7.1

(A7) *Independence: independent combination* (Harrington). For every pair of idealized preparations, one of each system, the two orders of idealization assign the same value to C; the resulting function on pairs of states, the independent product evaluated at C, is then the only extension of the expectation that is weak* continuous in each state separately, and it is symmetric in the two systems. Theorem 8.2

Several of the equivalences in Theorem A have nearly identical proofs; for instance, (A4) and (A7) express the same identity, through the Arens products, which are the two orders of idealization written as products of states, and through Harrington's lemma, respectively (see Section 8.3). The interest of these equivalences lies in their different meanings.

Theorem B (uniform stability in five languages). Let H_1 and H_2 be the Hilbert spaces of Alice's and Bob's systems, and let C be a bounded self-adjoint observable on $H_1 \otimes H_2$, with expectation $\varphi_C(\rho, \sigma) = \mathrm{Tr}\big((\rho \otimes \sigma)C\big)$ for $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$. Then the following are equivalent:

(B1) *Analysis: uniform order of idealization.* The observable C is uniformly stable: at every margin, its ladders, finite staircases like that of the order projection, have bounded length. Definition 9.5

(B2) *Combinatorics: decision trees* (Shelah). At every margin, the binary decision trees encoded by C have bounded height. Theorem 9.6

(B3) *Ultraproducts: families of experiments.* For every (equivalently, for some) nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, the ultrapower expectation $\Phi_C^{\mathcal{U}}(x, y) =$

$\lim_{k,\mathcal{U}} \varphi_C(x_k, y_k)$, that is, the expectation of C on families of experiments $x = (x_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $y = (y_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$, is stable. Theorem 9.6

(B4) *Growth: discrimination* (Sauer–Shelah). For every margin $\delta > 0$, the largest number of preparations that n adaptive probes can tell apart at margin δ, as a function of the positive integer n, is bounded by a polynomial, instead of being 2^n. Theorem 9.9

(B5) *Learnability: learning* (Littlestone). For every margin, in the threshold game of Section 9.5, the side of a fixed threshold on which the response of an unknown preparation lies can be predicted online with a number of mistakes bounded independently of the threshold. Proposition 9.13

Uniform stability implies stability, but not conversely (see Theorem 9.6 and Example 9.8). If there are maps u from $\mathfrak{D}(H_1)$ and v from $\mathfrak{D}(H_2)$ into a Hilbert space such that $\|u(\rho)\| \le 1$, $\|v(\sigma)\| \le 1$ and $\varphi_C(\rho, \sigma) = \langle u(\rho), v(\sigma)\rangle$ for all $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$, then, at every margin $\delta > 0$ and in every dimension, the number in (B4) is bounded by a polynomial in n of degree at most $16/\delta^2$, and the mistake bound in (B5) is at most $16/\delta^2$ (see Theorem 9.11).

Polynomial approximation is affine. In our work on deep computations [ADIW24, DIMW+26], which are idealized limits of computations of unbounded depth, tame behavior corresponds, roughly, to approximation by polynomials in finitely many realized values, and wild behavior to exponential growth. For bipartite observables the approximating polynomials can be taken to have degree one, and the reason is convexity: since expectations are affine in the state, for finitely many preparations $\rho_1, \dots, \rho_m \in \mathfrak{D}(H_1)$, weights $t_1, \dots, t_m \ge 0$ with $t_1 + \cdots + t_m = 1$ and every $\sigma \in \mathfrak{D}(H_2)$, the convex combination $\sum_{j=1}^m t_j \varphi_C(\rho_j, \sigma) = \varphi_C(\sum_{j=1}^m t_j \rho_j, \sigma)$ of actual responses is itself the response of a physical mixture, so the actual responses form a convex set. For a stable observable, idealized responses, which like actual ones are given by effective observables on Bob's system, lie in the closure of this set in the weak topology of the Banach space $B(H_2)$, and Mazur's theorem, by which a convex set has the same closure in the weak and in the norm topology, turns weak approximation by a convex set into uniform approximation by its elements. Thus idealized responses are uniform limits of *affine* combinations of actual responses (see Theorem 6.1 and Remark 6.4). In Theorems 9.11 and 6.5 we show that, for the swap test, Hong–Ou–Mandel coincidences and equality tests, two further bounds come from linear arguments: Novikoff's perceptron bound, a mistake bound for a linear classifier, limits the heights of decision trees, and Maurey's sampling argument, which replaces a convex combination by the average of a few random draws, limits the size of the simulating mixtures, both independently of the dimension.

Classical results and new results. For bounded functions on products of sets and for bilinear forms, the analogues of the equivalences of Theorem A are classical, as are the passage from ladders to trees, the characterization of uniform stability through ultrapowers, the Sauer–Shelah bound for trees, Littlestone's mistake bound, Novikoff's perceptron bound and Maurey's sampling argument; we cite them in the historical remarks at the end of the sections. What is new here is their formulation for bipartite observables through the two orders of idealization, the physical meaning of each, and the following results, which we have not found in the literature in this form: the criterion that pure probes, those given by unit vectors, suffice (see Proposition 5.2); the fact that countably additive randomization

never detects instability (see Proposition 7.2); the form of simulation by mixtures and its failure in the other order (see Theorem 6.1 and Example 6.2); the Monte Carlo bound on the size of simulating mixtures (see Theorem 6.5); the independent product of idealized preparations (see Theorem 8.2); the strict gap between stability and uniform stability for observables (see Example 9.8); the dimension-free tree bound for Hilbert-factorable observables (see Theorem 9.11); and the failure of stability under meets, the quantum-logical "and" (see Proposition 10.3). Several of these are short consequences of the classical results, and we present them as such.

Organization of the paper. Most sections are titled after the area they draw on. We begin with physics, where we make idealized preparations precise and introduce the two orders of idealization (see Section 2), and with analysis, where we develop the double limit property (see Section 3). The experiments listed above are treated in Section 4. The proof of Theorem A occupies Sections 5–8: topology, geometry and algebra (idealized responses, factorization and idealized states) in Section 5, convexity (simulation) in Section 6, probability (randomization) in Section 7, and independence (independent combination) in Section 8. We prove Theorem B through combinatorics in Section 9. With both dichotomies established, we turn to the order structure of the stable projections in Section 10, where we show that stability survives linear combinations, adjoints, norm limits and orthogonal sums, but not meets or strong limits. The paper ends with an outlook in Section 11 and open questions in Section 12.

2 Physics: two orders of idealization

We first give a precise meaning to an idealized preparation and show that two idealized preparations can be combined in two different orders. In the rest of the paper we compare these two orders.

2.1 Actual states and effective observables

As in the Introduction, the complex separable Hilbert spaces H_1 and H_2 describe Alice's and Bob's systems; throughout, the letter H denotes a complex separable Hilbert space, and we write $\mathfrak{D}(H)$ for the set of density operators on H, the positive operators of trace one. We follow the mathematicians' convention that the inner product $\langle \psi, \chi \rangle$ is linear in ψ, and write matrix elements in Dirac notation: for vectors ψ and χ of H and a bounded operator A on H, the bracket $\langle \chi | A | \psi \rangle$ stands for $\langle A\psi, \chi \rangle$, and the symbol $|\psi\rangle\langle\chi|$ denotes the rank-one operator $\xi \mapsto \langle \xi, \chi \rangle \psi$. The space $\mathfrak{T}(H)$ of trace-class operators on H consists of the operators τ whose trace norm $\|\tau\|_1 = \operatorname{Tr}|\tau|$, where $|\tau| = (\tau^*\tau)^{1/2}$, is finite. With this norm, the space $\mathfrak{T}(H)$ has as its dual the algebra $B(H)$ of bounded operators on H, each operator $A \in B(H)$ acting by $\tau \mapsto \operatorname{Tr}(\tau A)$. For a unit vector ψ, the pure state $|\psi\rangle\langle\psi|$, the orthogonal projection onto the line through ψ, is often written simply ψ.

Definition 2.1 (Bilinear form and slice maps). Let $C \in B(H_1 \otimes H_2)$ be given.

(1) The *bilinear form* of C is the map $m_C(\tau, \tau') = \operatorname{Tr}\big((\tau \otimes \tau')C\big)$ on $\mathfrak{T}(H_1) \times \mathfrak{T}(H_2)$.

(2) The *slice maps* of C are the operators $T_C \colon \mathfrak{T}(H_1) \to B(H_2)$ and $S_C \colon \mathfrak{T}(H_2) \to B(H_1)$ determined by

$$\operatorname{Tr}(\tau' T_C(\tau)) = m_C(\tau, \tau') = \operatorname{Tr}(\tau S_C(\tau')) \qquad \text{for all } \tau \in \mathfrak{T}(H_1) \text{ and } \tau' \in \mathfrak{T}(H_2).$$

The restriction of m_C to $\mathfrak{D}(H_1) \times \mathfrak{D}(H_2)$ is the expectation $\varphi_C(\rho, \sigma) = \mathrm{Tr}((\rho \otimes \sigma)C)$ of the Introduction, and both slice maps have norm at most $\|C\|$. In physical terms, for a density operator ρ on H_1, the operator $T_C(\rho) = \mathrm{Tr}_1((\rho \otimes I)C)$ is obtained from $(\rho \otimes I)C$ by the partial trace over Alice's system, the linear operation that sends $\tau \otimes B$ to $\mathrm{Tr}(\tau)\, B$ for trace-class τ on H_1 and bounded B on H_2: once Alice has prepared ρ, the observable C acts on Bob's system as the *effective observable* $T_C(\rho)$. In orthonormal bases (e_i) of H_1 and (f_k) of H_2, with i and j running over the index set of the first and k and l over that of the second, and for a unit vector $\psi = \sum_i \psi_i e_i$ of H_1, we have

$$\langle f_k | T_C(\psi) | f_l \rangle = \langle \psi \otimes f_k | C | \psi \otimes f_l \rangle = \sum_{i,j} \overline{\psi_i}\, \psi_j \, \langle e_i \otimes f_k | C | e_j \otimes f_l \rangle.$$

For bounded operators A on H_1 and B on H_2, the product $C = A \otimes B$ yields $T_C(\rho) = \mathrm{Tr}(\rho A)\, B$: every preparation of Alice's system leaves a multiple of B on Bob's side, and the slice map T_C has rank at most one.

2.2 Idealized states

A *state* of $B(H)$ is a positive linear functional ω with $\omega(I) = 1$. The state space $\mathcal{S}(B(H)) \subseteq B(H)^*$ carries the weak* topology, the topology of pointwise convergence on $B(H)$, in which it is compact. Density operators give the *normal*, or actual, states $A \mapsto \mathrm{Tr}(\rho A)$; we call the remaining states *idealized*. For instance, let $(e_k)_{k \in \mathbb{N}}$ be an orthonormal sequence in H (the energy eigenstates of an oscillator, say), and let $\mathcal{U}$ be a nonprincipal ultrafilter on $\mathbb{N}$; recall from the Introduction that every bounded sequence $(a_k)_{k \in \mathbb{N}}$ of complex numbers has a unique ultralimit $\lim_{k,\mathcal{U}} a_k$ along $\mathcal{U}$, the number a with $\{k : |a_k - a| < \varepsilon\} \in \mathcal{U}$ for every $\varepsilon > 0$. Then the state ω of $B(H)$ defined by

$$\omega(A) = \lim_{k,\mathcal{U}} \langle A e_k, e_k \rangle$$

vanishes on every compact operator, and it is therefore idealized, since every normal state $A \mapsto \mathrm{Tr}(\rho A)$ takes a positive value at the projection onto an eigenvector of ρ with a positive eigenvalue. When the vectors e_k are the energy eigenstates of an oscillator, the state ω assigns probability zero to every finite set of energy levels and describes an oscillator excited to infinite energy. More generally, for a sequence $(\rho_k)_{k \in \mathbb{N}}$ in $\mathfrak{D}(H)$, the *ultralimit state* $\lim_{k,\mathcal{U}} \rho_k$ is the state $A \mapsto \lim_{k,\mathcal{U}} \mathrm{Tr}(\rho_k A)$.

Lemma 2.2. *Let H be a complex separable Hilbert space. The normal states are weak* dense in the state space $\mathcal{S}(B(H))$. Every $\omega \in B(H)^*$ is a combination $\omega = c_1\omega_1 - c_2\omega_2 + ic_3\omega_3 - ic_4\omega_4$ of four states $\omega_1, \ldots, \omega_4$ with $0 \le c_j \le \|\omega\|$ for $j = 1, \ldots, 4$.*

Proof. For the first statement, suppose that a state ω is not in the weak* closure of the convex set $\mathfrak{D}(H)$. By the Hahn–Banach theorem, there is a self-adjoint $A \in B(H)$ with $\omega(A) > \sup_{\rho \in \mathfrak{D}(H)} \mathrm{Tr}(\rho A) = \sup \operatorname{spec} A$, which is impossible for a state. The second statement is the Jordan decomposition of the hermitian parts of ω (see [KR83, Section 4.3]). □

2.3 The two orders of idealization

Suppose that Alice's preparation is idealized to a state ω_1 of $B(H_1)$ and Bob's to a state ω_2 of $B(H_2)$. There are two natural ways of assigning an expectation to an observable $C \in B(H_1 \otimes H_2)$.

Alice idealized first. Each actual state σ of Bob's system leaves on Alice's side the effective observable $S_C(\sigma)$, to which Alice's idealized preparation assigns the expectation

$\omega_1(S_C\sigma)$. The map $\sigma \mapsto \omega_1(S_C\sigma)$ is a bounded linear functional on $\mathfrak{T}(H_2)$, so there is a unique operator $Z_{\omega_1} \in B(H_2)$ with

$$\operatorname{Tr}(\sigma Z_{\omega_1}) = \omega_1(S_C\sigma) \qquad \text{for all } \sigma \in \mathfrak{T}(H_2).$$

In other words, Z_{ω_1} is the image $S_C^*(\omega_1)$ of ω_1 under the transpose $S_C^*\colon B(H_1)^* \to B(H_2)$ of S_C, the map that composes each functional on $B(H_1)$ with S_C. The operator Z_{ω_1} is the effective observable that Alice's idealized preparation leaves on Bob's side. Bob then evaluates his state on it.

Bob idealized first. Symmetrically, Bob's idealized preparation leaves on Alice's side the operator $W_{\omega_2} = T_C^*(\omega_2)$ in $B(H_1)$, which satisfies $\operatorname{Tr}(\rho W_{\omega_2}) = \omega_2(T_C\rho)$ for all $\rho \in \mathfrak{T}(H_1)$, and Alice evaluates her state on it.

Definition 2.3 (The two products). For states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, and for $C \in B(H_1 \otimes H_2)$, we define

$$(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \omega_2(Z_{\omega_1}), \qquad (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) = \omega_1(W_{\omega_2}),$$

where $Z_{\omega_1} = S_C^*(\omega_1) \in B(H_2)$ and $W_{\omega_2} = T_C^*(\omega_2) \in B(H_1)$ are the images of ω_1 and ω_2 under the transposes of the slice maps of C, the effective observables that the two idealized preparations leave on Bob's and on Alice's side, respectively. The subscript names the system that is idealized first.

The next lemma explains the names: each product is an iterated limit, taken in one of the two orders.

Lemma 2.4. *Let $(\rho_i)_{i\in I}$ and $(\sigma_j)_{j\in J}$ be nets of normal states of $B(H_1)$ and of $B(H_2)$, indexed by directed sets I and J, with weak* limits ω_1 and ω_2, and let $C \in B(H_1 \otimes H_2)$ be given. Then we have*

$$(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \lim_j \lim_i \varphi_C(\rho_i, \sigma_j), \qquad (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) = \lim_i \lim_j \varphi_C(\rho_i, \sigma_j).$$

Moreover, the two products can be written through the bitransposes $T_C^{**}\colon B(H_1)^* \to B(H_2)^{**}$ and $S_C^{**}\colon B(H_2)^* \to B(H_1)^{**}$ of the slice maps, that is, the transposes of their transposes T_C^* and S_C^*:

$$(\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) = \langle T_C^{**}(\omega_1), \omega_2\rangle, \qquad (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \langle S_C^{**}(\omega_2), \omega_1\rangle.$$

In addition, the map $\omega_2 \mapsto (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$, defined on all states ω_2 of $B(H_2)$, is weak continuous, being evaluation at the operator $Z_{\omega_1} \in B(H_2)$ of Definition 2.3.*

Proof. For a normal state σ, we have

$$\operatorname{Tr}(\sigma Z_{\omega_1}) = \omega_1(S_C\sigma) = \lim_i \operatorname{Tr}(\rho_i S_C\sigma) = \lim_i \varphi_C(\rho_i, \sigma),$$

and then $\omega_2(Z_{\omega_1}) = \lim_j \operatorname{Tr}(\sigma_j Z_{\omega_1})$. The formula for $\otimes_{\mathsf{B}}$ is symmetric. Finally, the pairing $\langle T_C^{**}\omega_1, \omega_2\rangle$ equals $\omega_1(T_C^*\omega_2) = \omega_1(W_{\omega_2})$ by the definition of the bitranspose, and symmetrically for S_C^{**}. □

The bitranspose formulas of Lemma 2.4 make sense for all $\omega_1 \in B(H_1)^*$ and $\omega_2 \in B(H_2)^*$ and are bilinear in them; so extended, the two products are the two *Arens extensions* of the bilinear form m_C to the biduals $B(H_1)^*$ of $\mathfrak{T}(H_1)$ and $B(H_2)^*$ of $\mathfrak{T}(H_2)$ [Are51], and *Arens regularity* of m_C is the statement that they agree. In Corollary 2.5 below, we show that when one of the two states ω_1 and ω_2 is actual, the two products always agree.

Corollary 2.5 (Instability requires two idealized factors). *If ω_1 and ω_2 are states of $B(H_1)$ and of $B(H_2)$, respectively, and one of them is normal, then $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C)$ for every $C \in B(H_1 \otimes H_2)$.*

Proof. If ω_1 is given by a density operator ρ, the operator Z_{ω_1} is $T_C(\rho)$, and the product $(\omega_1 \otimes_{\mathsf{B}} \omega_2)(C)$ is $\operatorname{Tr}(\rho\, T_C^* \omega_2) = \omega_2(T_C \rho)$; both products therefore equal $\omega_2(T_C\rho)$. If ω_2 is given by a density operator σ, the operator $T_C^*(\sigma)$ coincides with $S_C(\sigma)$, and both products equal $\omega_1(S_C\sigma)$. □

3 Analysis: stability is the double limit property

In this section we develop the double limit property for arbitrary bounded functions on products: its ultrafilter form, the ladders that witness its failure, and its preservation under the operations used later.

Definition 3.1 (Double limit property). Let X and Y be sets. A bounded function $f\colon X \times Y \to \mathbb{C}$ is *stable* if for all sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y, we have

$$\lim_{l\to\infty} \lim_{k\to\infty} f(x_k, y_l) = \lim_{k\to\infty} \lim_{l\to\infty} f(x_k, y_l)$$

whenever both iterated limits exist. An observable $C \in B(H_1 \otimes H_2)$, where H_1 and H_2 are the Hilbert spaces of Alice's and Bob's systems, is *stable* if its expectation $\varphi_C(\rho, \sigma) = \operatorname{Tr}\big((\rho \otimes \sigma)C\big)$ is stable on the product $\mathfrak{D}(H_1) \times \mathfrak{D}(H_2)$ of the sets of density operators on H_1 and on H_2.

Two elementary facts are used throughout. First, every bounded array $(a_{kl})_{k,l\in\mathbb{N}}$ of complex numbers, with rows indexed by k and columns by l, has subsequences of rows and columns along which both iterated limits exist, and for a finite family of such arrays these subsequences can be chosen common to all arrays of the family: pass to a subsequence of columns along which every row converges (a diagonal argument), then to a subsequence of rows along which every remaining column converges, and then to further subsequences along which the two outer limits exist. Second, passing to subsequences does not change iterated limits that already exist.

The basic tool is the following extraction lemma, which turns a gap between two iterated ultrafilter limits into a *ladder*: an array that is close to one limit on and above the diagonal and close to the other below it. Recall that, for a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, the ultralimit $\lim_{k,\mathcal{U}} a_k$ of a bounded sequence $(a_k)_{k\in\mathbb{N}}$ of complex numbers is the unique number a with $\{k : |a_k - a| < \varepsilon\} \in \mathcal{U}$ for every $\varepsilon > 0$.

Lemma 3.2 (Ladder extraction). *Let $(a_{kl})_{k,l\in\mathbb{N}}$ be a bounded complex array, let $\mathcal{U}$ and $\mathcal{V}$ be nonprincipal ultrafilters on $\mathbb{N}$, and let*

$$L_1 = \lim_{k,\mathcal{U}} \lim_{l,\mathcal{V}} a_{kl}, \qquad L_2 = \lim_{l,\mathcal{V}} \lim_{k,\mathcal{U}} a_{kl}$$

be the two iterated limits. For every $\varepsilon > 0$ there are indices $k_1 < k_2 < \cdots$ and $l_1 < l_2 < \cdots$ such that $|a_{k_i l_m} - L_1| \le \varepsilon$ for $i \le m$ and $|a_{k_i l_m} - L_2| \le \varepsilon$ for $i > m$.

Proof. The indices are chosen alternately, each in a set that belongs to the relevant ultrafilter. Put $\alpha_k = \lim_{l,\mathcal{V}} a_{kl}$ and $\beta_l = \lim_{k,\mathcal{U}} a_{kl}$. Given $k_1, l_1, \ldots, k_{m-1}, l_{m-1}$, choose $k_m > k_{m-1}$ in the set

$$\{k :\ |\alpha_k - L_1| \le \varepsilon/2 \text{ and } |a_{kl_i} - \beta_{l_i}| \le \varepsilon/2 \text{ for all } i < m\},$$

which belongs to $\mathcal{U}$, and then $l_m > l_{m-1}$ in the set

$$\{l : \ |\beta_l - L_2| \le \varepsilon/2 \text{ and } |a_{k_i l} - \alpha_{k_i}| \le \varepsilon/2 \text{ for all } i \le m\},$$

which belongs to $\mathcal{V}$. For $i \le m$, we have $|a_{k_i l_m} - L_1| \le |a_{k_i l_m} - \alpha_{k_i}| + |\alpha_{k_i} - L_1| \le \varepsilon$, and for $i > m$, the choice of k_i after l_m ensures that $|a_{k_i l_m} - L_2| \le |a_{k_i l_m} - \beta_{l_m}| + |\beta_{l_m} - L_2| \le \varepsilon$. □

Corollary 3.3. *Let X and Y be sets, and let $f\colon X \times Y \to \mathbb{C}$ be a bounded function. Then the function f is stable if and only if, for all sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y and all nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$, we have*

$$\lim_{k,\mathcal{U}} \lim_{l,\mathcal{V}} f(x_k, y_l) = \lim_{l,\mathcal{V}} \lim_{k,\mathcal{U}} f(x_k, y_l).$$

Proof. If ordinary iterated limits exist, they equal the ultrafilter ones; this gives one direction. For the other, suppose that, for some sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y and some nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$, the two ultrafilter iterated limits L_1 and L_2 differ. Apply Lemma 3.2 to $a_{kl} = f(x_k, y_l)$ with $2\varepsilon < |L_1 - L_2|$, and pass to further subsequences $(k_{i_p})_{p\in\mathbb{N}}$ and $(l_{m_q})_{q\in\mathbb{N}}$ of the indices it provides, along which both ordinary iterated limits exist. For fixed p, we have $i_p \le m_q$ for all large q, so the limit over q is within ε of L_1, and hence so is the iterated limit $\lim_p \lim_q$. Symmetrically, the iterated limit $\lim_q \lim_p$ is within ε of L_2. The two iterated limits differ, so the function f is unstable. □

After the application of Lemma 3.2, the proof uses only the ladder inequalities, and it therefore shows the following: if f is a bounded function on $X \times Y$ and there are sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y, indices $k_1 < k_2 < \cdots$ and $l_1 < l_2 < \cdots$, complex numbers L_1 and L_2, and a real number ε with $2\varepsilon < |L_1 - L_2|$ such that the values $a_{k_i l_m} = f(x_{k_i}, y_{l_m})$ satisfy $|a_{k_i l_m} - L_1| \le \varepsilon$ for $i \le m$ and $|a_{k_i l_m} - L_2| \le \varepsilon$ for $i > m$, then the function f is unstable.

Lemma 3.4 (Preservation of stability). *Let X and Y be sets.*

(a) (Sums with one-variable coefficients.) *For every positive integer r, if $f_1, \dots, f_r$ are stable functions on $X \times Y$ and the functions $c_1, \dots, c_r\colon X \to \mathbb{C}$ and $d_1, \dots, d_r\colon Y \to \mathbb{C}$ are bounded, then the function $(x, y) \mapsto \sum_{j=1}^r c_j(x) d_j(y) f_j(x, y)$ is stable. In particular, linear combinations of stable functions and separated sums $\sum_{j=1}^r c_j(x) d_j(y)$ are stable.*

(b) (Uniform limits.) *If f and f' are bounded functions on $X \times Y$, the function f' is stable, and a real number $\varepsilon \ge 0$ satisfies $\sup_{X\times Y} |f - f'| \le \varepsilon$, then, for all sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y along which both iterated limits of f exist, these two limits differ by at most 2ε. In particular, uniform limits of stable functions are stable.*

(c) (Composition.) *If f is stable on $X \times Y$, then for all sets X' and Y' and all maps $g\colon X' \to X$ and $h\colon Y' \to Y$, the function $(x', y') \mapsto f(g(x'), h(y'))$ is stable on $X' \times Y'$.*

(d) (Reduction to states.) *For $C \in B(H_1 \otimes H_2)$, the expectation φ_C is stable on $\mathfrak{D}(H_1) \times \mathfrak{D}(H_2)$ if and only if the bilinear form $m_C(\tau, \tau') = \mathrm{Tr}\big((\tau \otimes \tau')C\big)$ of Definition 2.1 is stable on the product of the unit balls of $\mathcal{T}(H_1)$ and $\mathcal{T}(H_2)$.*

Proof. (a) Put $F = \sum_j c_j d_j f_j$, and let $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y be sequences along which both iterated limits of F exist. Pass to subsequences along which, for every j, both iterated limits of f_j exist and the numbers $c_j(x_k)$ and $d_j(y_l)$ converge to limits c_j^∞ and d_j^∞, respectively. By stability, the two iterated limits of f_j have a common value Λ_j, and both iterated limits of F along the subsequences equal $\sum_j c_j^\infty d_j^\infty \Lambda_j$.

(b) Let L_1 and L_2 be the two iterated limits of f along sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y along which both exist. Pass to subsequences along which both iterated limits of f' exist. They have a common value Λ, and both $|L_1 - \Lambda|$ and $|L_2 - \Lambda|$ are at most ε.

(c) Sequences in X' and Y' are mapped to sequences in X and Y.

(d) One direction is trivial. For the other, decompose each τ in the unit ball of $\mathcal{T}(H)$, where H is H_1 or H_2, as $\tau = c_1\rho_1 - c_2\rho_2 + ic_3\rho_3 - ic_4\rho_4$, using the positive and negative parts of the hermitian parts of τ, with $\rho_j \in \mathfrak{D}(H)$ and $0 \le c_j \le 1$ (choosing a fixed state when a part vanishes). These choices define maps $\tau \mapsto \rho_j(\tau)$ and bounded coefficients $\tau \mapsto c_j(\tau)$. For τ in the unit ball of $\mathcal{T}(H_1)$ and τ' in that of $\mathcal{T}(H_2)$, with $c'_{j'}$ and $\rho'_{j'}$ the coefficients and states obtained in the same way from τ', the value $m_C(\tau,\tau')$ is then a sum of sixteen terms $c_j(\tau)c'_{j'}(\tau')\varphi_C(\rho_j(\tau),\rho'_{j'}(\tau'))$, each multiplied by one of the factors ± 1 or $\pm i$, so parts (a) and (c) apply. □

3.1 Historical remarks

The double limit condition for a function on a product goes back to Grothendieck [Gro52], as a criterion for weak compactness in spaces of continuous functions. Independently, Shelah introduced the *order property* of a formula and defined stable theories as those in which no formula has it [She71]; a formula has the order property exactly when it has an infinite ladder in some model of the theory or, by compactness, ladders of every finite length. Krivine and Maurey introduced stable Banach spaces, those in which the function $(x,y) \mapsto \|x+y\|$ satisfies the double limit condition on bounded sets [KM81]. Iovino recognized that Shelah's order property is the two-valued case of the same condition, developed stability for real-valued structures on the basis of Krivine–Maurey stability [Iov94, Iov99a], and proved that a bounded function on a product of two sets is stable exactly when it has the form $(x,y) \mapsto \langle u(x), v(y)\rangle$ for bounded maps u into a reflexive Banach space and v into its dual [Iov99b]. Ben Yaacov and Usvyatsov [BYU10] developed local stability in the continuous logic of Ben Yaacov, Berenstein, Henson and Usvyatsov [BYBHU08], and Ben Yaacov derived definability of types directly from Grothendieck's theorem [BY14]. The extraction of ladders in Lemma 3.2 follows an argument of Iovino [Iov99b] (see also [Iov, Chapter 1]).

4 Experiments: the swap test, interference and equality tests are stable

We now decide stability for the observables listed in the Introduction. We show that in these examples, stability comes from one of two sources, approximation by finite sums of products or factorization through a Hilbert space (see Theorem 4.1 and Proposition 4.4), and instability from an encoding of the order of the natural numbers (see Theorem 4.7).

4.1 The minimal tensor product

We write $B(H_1) \odot B(H_2)$ for the space of finite sums of products $A \otimes B$ of bounded operators A on H_1 and B on H_2, the Hilbert spaces of Alice's and Bob's systems. Its norm closure is the minimal C*-tensor product $B(H_1) \otimes_{\min} B(H_2)$, which consists of the

bipartite observables that are uniform limits of finite sums of products of an observable of Alice's system with one of Bob's.

Theorem 4.1. *Every $C \in B(H_1) \otimes_{\min} B(H_2)$ is stable. More generally, let $C \in B(H_1 \otimes H_2)$ be given, with expectation $\varphi_C(\rho, \sigma) = \operatorname{Tr}((\rho \otimes \sigma)C)$. For all sequences $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $(\sigma_l)_{l\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$ along which both iterated limits of $\varphi_C(\rho_k, \sigma_l)$ exist, these two limits differ by at most $2\operatorname{dist}\big(C, B(H_1) \odot B(H_2)\big)$, the distance being taken in the operator norm.*

Proof. Finite sums of products are stable, and stability survives uniform approximation. Indeed, for a finite sum $C' = \sum_{j=1}^{r} A_j \otimes B_j$ of products, with $A_j \in B(H_1)$ and $B_j \in B(H_2)$, the expectation $\varphi_{C'}(\rho, \sigma) = \sum_j \operatorname{Tr}(\rho A_j) \operatorname{Tr}(\sigma B_j)$ is a separated sum, stable by Lemma 3.4(a), and $|\varphi_C - \varphi_{C'}| \le \|C - C'\|$; part (b) of the same lemma, applied with $\varepsilon = \|C - C'\|$, bounds the difference of the two iterated limits by $2\|C - C'\|$, and the infimum over all such C' completes the proof. □

Corollary 4.2.

(a) *Every compact operator on $H_1 \otimes H_2$, in particular every finite-rank projection, such as the projection onto a Bell state, a maximally entangled state of two qubits (quantum two-level systems), is stable.*

(b) *If H_1 or H_2 is finite-dimensional, then every $C \in B(H_1 \otimes H_2)$ is stable.*

Proof. (a) The rank-one operators $|\psi \otimes \chi\rangle\langle\psi' \otimes \chi'| = |\psi\rangle\langle\psi'| \otimes |\chi\rangle\langle\chi'|$, for $\psi, \psi' \in H_1$ and $\chi, \chi' \in H_2$, span a dense subspace of the compact operators, so Theorem 4.1 applies. (b) If H_1 has an orthonormal basis $e_1, \dots, e_n$, then the operator C equals $\sum_{i,j\le n} |e_i\rangle\langle e_j| \otimes C_{ij}$ for suitable $C_{ij} \in B(H_2)$, and therefore lies in $B(H_1) \odot B(H_2)$; the case of a finite-dimensional H_2 is symmetric. □

The slice map T_C of Definition 2.1, which sends a preparation of Alice's system to the effective observable that it leaves on Bob's side, gives a necessary condition for membership in the minimal tensor product.

Proposition 4.3. *If $C \in B(H_1) \otimes_{\min} B(H_2)$, then the slice map T_C is compact.*

Proof. The slice map $T_{A\otimes B} = \operatorname{Tr}(\cdot\, A)B$ has rank at most one, and $\|T_C - T_{C'}\| \le \|C - C'\|$ for every $C' \in B(H_1 \otimes H_2)$, in particular for the finite sums of products that approximate C; norm limits of finite-rank operators are compact. □

4.2 Factorization through a Hilbert space

In Section 4.3 and Proposition 4.6 we show that several physically basic observables lie outside the minimal tensor product and are nevertheless stable, because the expectation of each is the inner product of a vector that depends only on Alice's preparation with a vector that depends only on Bob's.

Proposition 4.4. *Let X and Y be sets, let $\mathcal{H}$ be a Hilbert space, and let $u\colon X \to \mathcal{H}$ and $v\colon Y \to \mathcal{H}$ be bounded maps. Then the function $f(x, y) = \langle u(x), v(y)\rangle$ is stable. In particular, if $C \in B(H_1 \otimes H_2)$ and there are bounded maps $u\colon \mathfrak{D}(H_1) \to \mathcal{H}$ and $v\colon \mathfrak{D}(H_2) \to \mathcal{H}$ such that the expectation $\varphi_C(\rho, \sigma) = \operatorname{Tr}\big((\rho \otimes \sigma)C\big)$ equals $\langle u(\rho), v(\sigma)\rangle$ for all $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$, then the observable C is stable.*

Proof. Bounded sets in a Hilbert space are relatively weakly sequentially compact. Given sequences $(x_k)_{k\in\mathbb{N}}$ in X and $(y_l)_{l\in\mathbb{N}}$ in Y along which both iterated limits exist, pass to subsequences $(x_{k_i})_{i\in\mathbb{N}}$ and $(y_{l_j})_{j\in\mathbb{N}}$ with $u(x_{k_i}) \rightharpoonup \xi$ and $v(y_{l_j}) \rightharpoonup \xi'$ weakly, for some vectors ξ and ξ' of $\mathcal{H}$. For each i, we have $\lim_j \langle u(x_{k_i}), v(y_{l_j})\rangle = \langle u(x_{k_i}), \xi'\rangle$, which tends to $\langle \xi, \xi'\rangle$; the other iterated limit is also $\langle \xi, \xi'\rangle$. □

4.3 The swap operator, symmetrization, and Hong–Ou–Mandel interference

Let K be a complex separable Hilbert space, which in this subsection is the Hilbert space of both systems. The *swap* $F \in B(K \otimes K)$ is the unitary operator with $F(\psi \otimes \chi) = \chi \otimes \psi$ for all $\psi, \chi \in K$. Since $F = F^*$ and $F^2 = I$, the operators

$$P_+ = \tfrac{1}{2}(I + F), \qquad P_- = \tfrac{1}{2}(I - F)$$

are the projections onto the symmetric and the antisymmetric subspaces of $K \otimes K$, the eigenspaces of F for the eigenvalues 1 and -1, respectively.

Lemma 4.5 (Swap trick). *For all $\tau_1, \tau_2 \in \mathcal{T}(K)$, we have* $\operatorname{Tr}((\tau_1 \otimes \tau_2)F) = \operatorname{Tr}(\tau_1\tau_2)$.

Proof. For $\tau_1 = |\psi\rangle\langle\psi'|$ and $\tau_2 = |\chi\rangle\langle\chi'|$, with $\psi, \psi', \chi, \chi' \in K$, one computes $(\tau_1 \otimes \tau_2)F = |\psi \otimes \chi\rangle\langle\chi' \otimes \psi'|$, whose trace is $\langle\psi, \chi'\rangle\langle\chi, \psi'\rangle = \operatorname{Tr}(\tau_1\tau_2)$. Extend by bilinearity and continuity. □

By Lemma 4.5, for density operators ρ and σ on K, the expectation of the swap is $\varphi_F(\rho, \sigma) = \operatorname{Tr}(\rho\sigma)$, the inner product $\langle\rho, \sigma\rangle = \operatorname{Tr}(\sigma^*\rho)$ of ρ and σ in the Hilbert space $S_2(K)$ of Hilbert–Schmidt operators, the operators τ on K with $\operatorname{Tr}(\tau^*\tau) < \infty$, whence

$$\varphi_{P_\pm}(\rho, \sigma) = \tfrac{1}{2}(1 \pm \operatorname{Tr}(\rho\sigma)).$$

The operators F and $P_\pm$ are therefore stable (see Proposition 4.4). On the other hand, the slice map $T_F(\rho) = \rho$ is the inclusion $\mathcal{T}(K) \to B(K)$, which is not compact when K is infinite-dimensional, since the pure states given by the vectors of an orthonormal sequence have images at mutual distance one; the slice maps $T_{P_\pm} = \frac{1}{2}(\operatorname{Tr}(\,\cdot\,)I \pm T_F)$ differ from $\pm\frac{1}{2}T_F$ by an operator of rank one, and are therefore not compact either. Hence, by Proposition 4.3, when K is infinite-dimensional, the operators F and $P_\pm$ do not belong to $B(K) \otimes_{\min} B(K)$.

The overlap $\operatorname{Tr}(\rho\sigma)$ appears in three experiments:

- *The swap test* [BCWdW01]. An ancilla, an auxiliary qubit (a quantum two-level system), controls a swap of two registers, the quantum systems prepared in ρ and σ. The test accepts with probability $\frac{1}{2}(1 + \operatorname{Tr}\rho\sigma) = \varphi_{P_+}(\rho, \sigma)$, whence the frequency of acceptance in repeated runs estimates the overlap $\operatorname{Tr}\rho\sigma$ of the two states.

- *Symmetric and antisymmetric subspaces.* For two independently prepared systems with the same Hilbert space, the number $\varphi_{P_\pm}(\rho, \sigma)$ is the probability that a projective test for exchange symmetry (respectively, antisymmetry) succeeds. For genuinely identical particles a product state $\rho \otimes \sigma$ is in general not admissible, since the joint state of two identical particles is supported on the symmetric subspace (for bosons) or on the antisymmetric subspace (for fermions), and by the formula for $\varphi_{P_\pm}$ above the state $\rho \otimes \sigma$ is so supported only for bosons and $\rho = \sigma$ pure. The interpretation of $\varphi_{P_\pm}(\rho, \sigma)$ as a probability therefore concerns two distinguishable particles, or the internal degrees of freedom of two identical ones, as in the next item.

- *Hong–Ou–Mandel interference* [HOM87]. Two single photons whose internal (spectral and polarization) states are ρ and σ enter the two input ports of a balanced (50:50) beam splitter, a device, such as a half-silvered mirror, that transmits each single incoming photon with probability $\frac{1}{2}$ and reflects it otherwise. The probability of a coincidence, a detection at both output ports, is $\frac{1}{2}(1 - \operatorname{Tr}\rho\sigma) = \varphi_{P_-}(\rho, \sigma)$, which vanishes for equal pure states: the two photons then leave by the same port (bunching).

4.4 Equality and order of discrete spectra

Let A and B be self-adjoint operators on H_1 and H_2, possibly unbounded, representing quantities of Alice's and of Bob's system, respectively, with pure point spectra, that is, with orthonormal bases of eigenvectors. Their spectral decompositions are $A = \sum_\lambda \lambda \Pi^A_\lambda$ and $B = \sum_\lambda \lambda \Pi^B_\lambda$, where λ runs over the reals and Π^A_λ and Π^B_λ are the projections onto the eigenspaces of A and of B for λ (zero unless λ is an eigenvalue). The events "equal values" and "Bob's value is larger" are described by the projections

$$[A = B] = \sum_\lambda \Pi^A_\lambda \otimes \Pi^B_\lambda, \qquad [A < B] = \sum_{\lambda<\lambda'} \Pi^A_\lambda \otimes \Pi^B_{\lambda'}.$$

Proposition 4.6 (Equality is stable, order is not). *Let A and B be self-adjoint operators on H_1 and H_2, possibly unbounded, with pure point spectra, and let $[A = B]$ and $[A < B]$ be the projections onto the events "equal values" and "Bob's value is larger" defined at the start of this subsection.*

(1) *The projection $[A = B]$ is stable. If A and B have infinitely many common eigenvalues, then the projection $[A = B]$ does not belong to $B(H_1) \otimes_{\min} B(H_2)$.*

(2) *If A or B has finitely many eigenvalues, then the projection $[A < B]$ is stable.*

(3) *If the point spectra of A and B are both unbounded above, then the projection $[A < B]$ is unstable.*

Proof. (1) We have $\varphi_{[A=B]}(\rho, \sigma) = \sum_\lambda \operatorname{Tr}(\rho\Pi^A_\lambda)\operatorname{Tr}(\sigma\Pi^B_\lambda) = \langle u(\rho), v(\sigma)\rangle_{\ell^2}$, where $u(\rho) = (\operatorname{Tr}\rho\Pi^A_\lambda)_\lambda$ and $v(\sigma) = (\operatorname{Tr}\sigma\Pi^B_\lambda)_\lambda$ are vectors of ℓ^2-norm at most one, and the projection is stable by Proposition 4.4. For each common eigenvalue λ, pick a unit vector ψ_λ in the range of Π^A_λ. The slice map $T_{[A=B]}$ sends ψ_λ to Π^B_λ, and these are mutually orthogonal nonzero projections, so the map is not compact; the conclusion now follows from Proposition 4.3.

(2) If A has finitely many eigenvalues, then the projection $[A < B] = \sum_\lambda \Pi^A_\lambda \otimes (\sum_{\lambda'>\lambda} \Pi^B_{\lambda'})$ is a finite sum of products, stable by Theorem 4.1. If B has finitely many eigenvalues, the same argument applies to $[A < B] = \sum_{\lambda'} (\sum_{\lambda<\lambda'} \Pi^A_\lambda) \otimes \Pi^B_{\lambda'}$.

(3) Since the point spectra of A and B are both unbounded above, we can choose real numbers $\lambda_1 < \lambda'_1 < \lambda_2 < \lambda'_2 < \cdots$ such that, for every $k \in \mathbb{N}$, the operators A and B have unit eigenvectors ψ_k and χ_k with eigenvalues λ_k and λ'_k, respectively. For $k, l \in \mathbb{N}$, the value $\varphi_{[A<B]}(\psi_k, \chi_l)$ is then 1 if $k \le l$ and 0 if $k > l$, so the iterated limit taken first in k is 0, and the one taken first in l is 1. □

4.5 The order projection

We return to the order projection of the Introduction, the sharpest example of instability: although it is diagonal in a product basis, and hence classical, its two iterated limits are as far apart as the expectations of a projection allow.

Theorem 4.7 (Instability of the order projection). *Let $(e_k)_{k\in\mathbb{N}}$ and $(f_l)_{l\in\mathbb{N}}$ be orthonormal sequences in H_1 and H_2, write e_k and f_l also for the pure states they define, and let $P_{k<l}$ be the projection onto the closed span of $\{e_k \otimes f_l : k < l\}$. Then, for all $k, l \in \mathbb{N}$, the expectation $\varphi_{P_{k<l}}(e_k, f_l) = \mathrm{Tr}\big((e_k \otimes f_l)P_{k<l}\big)$ is 1 if $k < l$ and 0 otherwise, so*

$$\lim_l \lim_k \varphi_{P_{k<l}}(e_k, f_l) = 0 \neq 1 = \lim_k \lim_l \varphi_{P_{k<l}}(e_k, f_l).$$

Hence the projection $P_{k<l}$ is unstable, and its distance to $B(H_1) \otimes_{\min} B(H_2)$ is at least $\frac{1}{2}$.

Proof. The values form the staircase

$$\begin{pmatrix} & f_1 & f_2 & f_3 & \cdots \\ e_1 & 0 & 1 & 1 & \cdots \\ e_2 & 0 & 0 & 1 & \cdots \\ e_3 & 0 & 0 & 0 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{pmatrix}.$$

Each column tends to 0 and each row to 1. The bound on the distance follows from Theorem 4.1. □

The table below summarizes the examples; uniform stability is introduced in Section 9.2 below, and the entries in its column are justified by Theorem 9.6, Proposition 9.7 and Example 9.8. Uniform stability corresponds to what is called superstability in Banach space theory, where a space is superstable when all its ultrapowers are stable [Ray83]; we prefer the name uniform stability because superstability has a different meaning in model theory [She71]. The last column of the table uses superreflexivity: a Banach space is *superreflexive* if every Banach space finitely representable in it, one whose finite-dimensional subspaces are almost isometric to subspaces of it, is reflexive; equivalently, all its ultrapowers are reflexive [Hei80].

Observable	Stable	Uniformly stable	Mechanism
$A \otimes B$, $A \otimes I + I \otimes B$	yes	yes	minimal tensor product
compact, Bell projections	yes	yes	minimal tensor product
any C, one factor finite-dimensional	yes	yes	finite sum of products
swap F; P_+ (swap test); P_- (Hong–Ou–Mandel)	yes	yes	factors through $S_2(K)$
$[A = B]$, pure point spectra	yes	yes	factors through ℓ^2
$[A < B]$, one spectrum finite	yes	yes	finite sum of products
sector comparison C_{sec} (see Example 9.8 below)	yes	**no**	reflexive, not superreflexive
$[A < B]$, point spectra unbounded above	**no**	**no**	encodes the order of $\mathbb{N}$
order projection $P_{k<l}$	**no**	**no**	encodes the order of $\mathbb{N}$

5 Topology: idealized responses are observables

Here we prove that an observable is stable if and only if the responses of every idealized preparation of Alice's system to all probes, idealized ones included and idealized first, are the expectations, in those probes, of a single observable of Bob's system. In topological terms, this says that the slice map of the observable, which sends each preparation of Alice's system to the effective observable that it leaves on Bob's side, maps bounded sets to relatively weakly compact sets. We then derive the geometric and algebraic forms of the same condition, and show that pure probes suffice to test it.

5.1 Topology: weak compactness

The proofs in this section and in Section 6 rest on five classical theorems about a Banach space X, which we recall in plain terms. A bounded operator is *weakly compact* if it maps the unit ball to a relatively weakly compact set.

- *Eberlein–Šmulian theorem* [Ebe47, Š40]. A subset of X is relatively weakly compact if and only if every sequence in it has a weak cluster point in X.
- *Gantmacher's theorem* [Gan40]. An operator $T\colon X \to Y$ between Banach spaces is weakly compact if and only if its bitranspose T^{**}, the transpose of its transpose, maps the bidual $X^{**} = (X^*)^*$ into Y, regarded as a subspace of Y^{**}.
- *Mazur's theorem* [Maz33]. The weak closure and the norm closure of a convex subset of X coincide.
- *Krein's theorem* [KŠ40]. The closed convex hull of a relatively weakly compact subset of X is weakly compact.
- *Davis–Figiel–Johnson–Pełczyński theorem* [DFJP74]. Every weakly compact operator factors through a reflexive Banach space.

Theorem 5.1 (Characterization of stability). *For every $C \in B(H_1 \otimes H_2)$, where H_1 and H_2 are the Hilbert spaces of Alice's and Bob's systems, the following are equivalent:*

(1) *The observable C is stable.*

(2) *The bilinear form m_C of Definition 2.1 is stable on the product of the unit balls of $\mathfrak{T}(H_1)$ and $\mathfrak{T}(H_2)$.*

(3) (Topology.) *The slice map $T_C \colon \mathfrak{T}(H_1) \to B(H_2)$ of Definition 2.1 is weakly compact.*

(4) (Geometry.) *There are a reflexive Banach space E and bounded operators $a \colon \mathfrak{T}(H_1) \to E$ and $J \colon E \to B(H_2)$ with $T_C = J \circ a$.*

(5) (Algebra.) *The two products of Definition 2.3, one for each order of idealization, agree at C: for all states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, we have $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C)$. Equivalently, the bitranspose T_C^{**} maps every state ω_1 of $B(H_1)$ into $B(H_2)$, namely to the effective observable Z_{ω_1} that ω_1 leaves on Bob's side (see Definition 2.3).*

(5′) *The equality in (5) holds for all ultralimit states $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$ and $\omega_2 = \lim_{l,\mathcal{V}} \sigma_l$ of sequences $(\rho_k)_{k\in\mathbb{N}}$ and $(\sigma_l)_{l\in\mathbb{N}}$ of normal states of $B(H_1)$ and of $B(H_2)$ along nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$.*

(6) (Separate continuity.) *There is a function $\hat{\varphi} \colon \mathcal{S}(B(H_1)) \times \mathcal{S}(B(H_2)) \to \mathbb{C}$, weak* continuous in each variable separately, that extends the expectation φ_C, in the sense that $\hat{\varphi}(\rho, \sigma) = \varphi_C(\rho, \sigma)$ for all normal states ρ and σ.*

We prove Theorem 5.1 in three steps, one for each language: in this subsection we show that (1) and (2) are equivalent and imply (3); in Section 5.2 we show that (3) and (4) are equivalent; and in Section 5.3 we close the circle through (5), (5′) and (6).

Proof of the first step. The equivalence of (1) and (2) is Lemma 3.4(d).

For the implication from (1) to (3), the idea is that the effective observable that an ultralimit of preparations leaves on Bob's side is a weak cluster point of their effective observables, and that any failure of this produces a ladder. Let $(\rho_k)_{k\in\mathbb{N}}$ be a sequence in $\mathfrak{D}(H_1)$, let $\mathcal{U}$ be a nonprincipal ultrafilter on $\mathbb{N}$, and let $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$. We claim that $\lim_{k,\mathcal{U}} \omega_2(T_C\rho_k) = \omega_2(Z_{\omega_1})$ for every state ω_2 of $B(H_2)$. Suppose instead that, for some state ω_2 of $B(H_2)$, the left side is L_1 and the right side is $L_2 \neq L_1$, and let $\varepsilon = |L_1 - L_2|/5$. By Lemma 2.4, we have $\lim_{k,\mathcal{U}} \operatorname{Tr}(\sigma\, T_C\rho_k) = \operatorname{Tr}(\sigma Z_{\omega_1})$ for every normal state σ. Choose recursively indices $k_1 < k_2 < \cdots$ and normal states $\sigma_1, \sigma_2, \ldots$ as follows. The index k_m is chosen in the set

$$\{k : \ |\omega_2(T_C\rho_k) - L_1| \le \varepsilon \text{ and } |\operatorname{Tr}(\sigma_i T_C\rho_k) - \operatorname{Tr}(\sigma_i Z_{\omega_1})| \le \varepsilon \text{ for all } i < m\},$$

which belongs to $\mathcal{U}$. The state σ_m is a normal state with $|\operatorname{Tr}(\sigma_m T_C\rho_{k_i}) - \omega_2(T_C\rho_{k_i})| \le \varepsilon$ for all $i \le m$ and $|\operatorname{Tr}(\sigma_m Z_{\omega_1}) - \omega_2(Z_{\omega_1})| \le \varepsilon$; such a state exists by Lemma 2.2. The value $\varphi_C(\rho_{k_i}, \sigma_m)$ is then within 2ε of L_1 for $i \le m$ and within 2ε of L_2 for $i > m$. Because $4\varepsilon < |L_1 - L_2|$, the paragraph after Corollary 3.3 shows that C is unstable, a contradiction. Since the claim extends from states to all of $B(H_2)^*$ by Lemma 2.2, the operator Z_{ω_1} is a weak cluster point of $(T_C\rho_k)$. The Eberlein–Šmulian theorem now shows that the set $T_C(\mathfrak{D}(H_1))$ is relatively weakly compact. Finally, the decomposition in the

proof of Lemma 3.4(d) places the image of the unit ball under T_C in a sum of four sets of the form $\{c\,A : |c| \le 1,\ A \in T_C(\mathfrak{D}(H_1))\}$, each of which is relatively weakly compact, so the slice map T_C is weakly compact. □

5.2 Geometry: factorization through a reflexive space

Condition (4) of Theorem 5.1 says that the expectation factors, by bounded affine maps, through the pairing of a reflexive space with its dual: with $u(\rho) = a(\rho)$ and $v(\sigma) = J^*(\sigma)$, the diagram

$$\begin{array}{ccc} \mathfrak{D}(H_1) \times \mathfrak{D}(H_2) & \xrightarrow{(u,v)} & E \times E^* \\ & \searrow^{\varphi_C} & \downarrow{\scriptstyle \langle\cdot,\cdot\rangle} \\ & & \mathbb{C} \end{array} \tag{2}$$

commutes.

Proof of the second step. We show that (3) and (4) are equivalent. If T_C is weakly compact, it factors through a reflexive space by the Davis–Figiel–Johnson–Pełczyński theorem, and the expectation becomes $\varphi_C(\rho,\sigma) = \mathrm{Tr}(\sigma\, J(a\rho)) = \langle a(\rho), J^*(\sigma)\rangle$, which is the commutativity of (2). Conversely, the image under a of the unit ball is bounded, hence relatively weakly compact in the reflexive space E, and J is weak-to-weak continuous, so the slice map $T_C = J \circ a$ is weakly compact. □

5.3 Algebra: the two orders of idealization agree

The algebraic form of stability is Arens regularity: the two products of Definition 2.3 agree. The proof of this equivalence closes the circle of implications in Theorem 5.1, and it also shows that the agreement of the two products is equivalent to the separate continuity in (6), which returns to topology.

Proof of the third step. We show that (3) and (5) are equivalent, that (5) implies (5′), which implies (1), and that (5) and (6) are equivalent; together with the first two steps, this proves the theorem. We begin with (3) and (5). For a state ω_1 and a normal state σ, we have $\langle T_C^{**}\omega_1, \sigma\rangle = \omega_1(T_C^*\sigma) = \omega_1(S_C\sigma) = \mathrm{Tr}(\sigma Z_{\omega_1})$. Hence the functional $T_C^{**}(\omega_1)$ belongs to $B(H_2)$ if and only if $T_C^{**}(\omega_1) = Z_{\omega_1}$, and by Lemma 2.4, this happens if and only if $(\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) = \omega_2(Z_{\omega_1}) = (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$ for all states ω_2. Since states span $B(H_1)^*$ (see Lemma 2.2), condition (5) says that $T_C^{**}(B(H_1)^*) \subseteq B(H_2)$, which is (3) by Gantmacher's theorem.

The implication from (5) to (5′) is trivial. For the implication from (5′) to (1), suppose that C is unstable. There are sequences in $\mathfrak{D}(H_1)$ and $\mathfrak{D}(H_2)$ along which the ordinary iterated limits of φ_C exist and differ, and by Lemma 2.4, the ultralimit states along any nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$ reproduce these limits as the two products, which therefore differ, contrary to (5′).

Finally, we turn to (5) and (6). By Lemma 2.4, the maps $\omega_2 \mapsto (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \omega_2(Z_{\omega_1})$ and $\omega_1 \mapsto (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) = \omega_1(W_{\omega_2})$, evaluations at the operators that the two idealized preparations leave on Bob's and on Alice's side, respectively, are weak* continuous, so under (5) their common value is separately continuous. Conversely, let $\hat{\varphi}$ be as in (6). For a normal state σ and a net (ρ_i) of normal states of $B(H_1)$, indexed by a directed set, with limit ω_1, continuity in the first variable gives $\hat{\varphi}(\omega_1,\sigma) = \lim_i \varphi_C(\rho_i,\sigma) = (\omega_1 \otimes_{\mathsf{A}} \sigma)(C)$. Both $\hat{\varphi}(\omega_1,\cdot\,)$ and $(\omega_1 \otimes_{\mathsf{A}} \cdot\,)(C)$ are weak* continuous and normal states are dense, so the

two functions agree: $\hat{\varphi}(\omega_1, \omega_2) = (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$ for all pairs of states. Symmetrically, the function $\hat{\varphi}$ also coincides with $(\omega_1, \omega_2) \mapsto (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C)$, and (5) follows. □

5.4 Pure probes suffice

Pure states, those given by unit vectors, are the extreme points of the set of density operators and the simplest states to describe. In the next proposition we show that pure preparations and pure probes detect every instability.

Proposition 5.2 (Pure probes suffice). *Let $C \in B(H_1 \otimes H_2)$ be given. If the restriction of the expectation $\varphi_C(\rho, \sigma) = \operatorname{Tr}((\rho \otimes \sigma)C)$ to pairs of pure states ρ of Alice's system and σ of Bob's is stable, then the observable C is stable.*

Proof. The idea is to represent effective observables as continuous functions on the closure of the pure states of $B(H_2)$, where Grothendieck's criterion needs to test only a dense set of points. For density operators we have $\varphi_{C^*} = \overline{\varphi_C}$, so the hermitian parts $\frac{1}{2}(C + C^*)$ and $\frac{1}{2i}(C - C^*)$ of C inherit the hypothesis, and we may assume that C is self-adjoint; then the operator $T_C(\rho)$ is self-adjoint for $\rho \geq 0$. Let $\Omega \subseteq \mathcal{S}(B(H_2))$ be the weak* closure of the pure states, a compact space. For a self-adjoint operator A, the function $\hat{A}(\omega) = \omega(A)$ on Ω satisfies $\sup_\Omega |\hat{A}| = \|A\|$, so the map $A \mapsto \hat{A}$ is a linear isometry of the self-adjoint part of $B(H_2)$ onto a closed subspace of the space $\mathcal{C}(\Omega)$ of continuous functions on Ω, with the supremum norm. Since the pure states are dense in Ω, restriction to the set P_2 of pure states of Bob's system is a linear isometry of $\mathcal{C}(\Omega)$ onto a closed subspace of the space $\ell^\infty(\mathsf{P}_2)$ of bounded functions on P_2. By Grothendieck's theorem [Gro52] (see also [Ptá63]), a bounded subset of $\ell^\infty(\mathsf{P}_2)$ is relatively weakly compact exactly when evaluation, as a function on the product of the subset with P_2, is stable in the sense of Definition 3.1. Applied to the double limit property of the hypothesis, it shows that the restrictions of the functions $\omega \mapsto \omega(T_C\psi)$, for pure states ψ of Alice's system, form a relatively weakly compact subset of $\ell^\infty(\mathsf{P}_2)$, and hence the functions themselves form a relatively weakly compact subset of $\mathcal{C}(\Omega)$. Relative weak compactness passes to the closed subspace by the Hahn–Banach theorem, so the set of operators $T_C\psi$, for ψ pure, is relatively weakly compact, and by Krein's theorem its closed convex hull is weakly compact. This hull contains $T_C(\mathfrak{D}(H_1))$, because every density operator is a trace-norm convergent convex combination of pure states. As at the end of the first step of the proof of Theorem 5.1, it follows that the slice map T_C is weakly compact, whence the observable C is stable by Theorem 5.1. □

5.5 Historical remarks

For general bounded bilinear forms, the equivalence of (2) and (5) in Theorem 5.1 is Pym's iterated limit criterion for Arens regularity [Pym65], and the equivalence of (3) and (5) is due to Ülger [Ü87]; the Arens products themselves were introduced by Arens [Are51]. The equivalence of (1) and (3) is an instance of Grothendieck's theorem [Gro52], which equates the double limit condition with relative weak compactness; the proof given here is a self-contained version adapted to states. The factorization (4) is the operator form of the theorem of Iovino that a bounded function on a product of two sets is stable exactly when it has the form $(x, y) \mapsto \langle u(x), v(y) \rangle$ for bounded maps u into a reflexive space and v into its dual [Iov99b, Theorem 3.3], whose proof combines the Krivine–Maurey and Raynaud constructions [KM81, Ray83] with the Davis–Figiel–Johnson–Pełczyński factorization [DFJP74].

6 Convexity: finite mixtures simulate idealized preparations

Idealized preparations are not available in the laboratory, but finite mixtures of actual ones are: Alice chooses one of finitely many actual preparations at random, with prescribed probabilities, and prepares it. Stability, we show here, is exactly the possibility of reproducing to arbitrary accuracy the responses of an idealized preparation of Alice's system by such mixtures, uniformly over all probes, that is, over all preparations of Bob's system, idealized ones included, provided that Bob's system, the probing one, is idealized first. For Hilbert-factorable observables, those whose expectation is the inner product of two vectors of a Hilbert space that depend boundedly on the two preparations, we also bound the size of the mixtures.

Theorem 6.1 (Affine simulation). *Let $C \in B(H_1 \otimes H_2)$ be given, where H_1 and H_2 are the Hilbert spaces of Alice's and Bob's systems, let T_C be the slice map of Definition 2.1, which sends a density operator ρ on H_1 to the effective observable $T_C(\rho)$ that it leaves on Bob's side, and let $\otimes_{\mathsf{A}}$ and $\otimes_{\mathsf{B}}$ be the products of Definition 2.3, in which Alice's and Bob's system, respectively, is idealized first. Then the following are equivalent:*

(1) *The observable C is stable.*

(2) (Bob idealized first.) *For every state ω_1 of $B(H_1)$ and every $\varepsilon > 0$, there is $\rho \in \mathfrak{D}(H_1)$ such that for all states ω_2 of $B(H_2)$, we have*

$$\big|(\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) - \omega_2(T_C(\rho))\big| \leq \varepsilon.$$

(3) (Tail mixtures.) *For every sequence $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$, every nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, every $n \in \mathbb{N}$ and every $\varepsilon > 0$, there is a finite convex combination $\rho = \sum_{k\geq n} t_k \rho_k$ such that for all states ω_2 of $B(H_2)$, we have*

$$\big|(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) - \omega_2(T_C(\rho))\big| \leq \varepsilon,$$

where $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$ is the ultralimit state of the sequence along $\mathcal{U}$.

If these conditions hold, then the two products agree at C, and the density operator ρ in (2) can be taken to be a finite mixture of elements of any given tail of any net in $\mathfrak{D}(H_1)$ that converges to ω_1 in the weak topology.*

Proof. We first show that (1) implies (2) and (3). The idea is that, under stability, the idealized response is an actual operator in the weak closure of a convex set of actual responses, and Mazur's theorem converts weak approximation into norm approximation. By Theorem 5.1, the bitranspose T_C^{**} of the slice map sends ω_1 to the operator $Z_{\omega_1} \in B(H_2)$, the effective observable that Alice's idealized preparation leaves on Bob's side (see Section 2), and both products equal $\omega_2(Z_{\omega_1})$. Let $\varepsilon > 0$ be given. If $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$ for a sequence $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, then, for every $n \in \mathbb{N}$, the state ω_1 lies in the weak* closure of the convex hull of $\{\rho_k : k \geq n\}$. The bitranspose T_C^{**} is weak*-to-weak* continuous, and on $B(H_2)$ the weak* topology of $B(H_2)^{**}$ is the weak topology. Consequently the operator Z_{ω_1} lies in the weak closure of the convex set $T_C(\mathrm{conv}\{\rho_k : k \geq n\})$, and therefore, by Mazur's theorem, in its norm closure: there is ρ in that convex hull with $\|Z_{\omega_1} - T_C(\rho)\| \leq \varepsilon$, and then $|\omega_2(Z_{\omega_1} - T_C\rho)| \leq \varepsilon$ for every state ω_2. The same argument applies to the tails of every net in $\mathfrak{D}(H_1)$ that converges to ω_1 in

the weak* topology; since every state ω_1 is the limit of such a net (see Lemma 2.2), we obtain (2) and the final assertion of the theorem.

Proof that (2) implies (1). By Lemma 2.4, the functional $G = T_C^{**}(\omega_1) - T_C(\rho)$ on $B(H_2)^*$ satisfies $|G(\omega_2)| \leq \varepsilon$ for all states ω_2, and Lemma 2.2 bounds its norm by 4ε. Since ε is arbitrary and $B(H_2)$ is norm closed in its bidual, the functional $T_C^{**}(\omega_1)$ lies in $B(H_2)$, and this holds for every state ω_1. By Theorem 5.1(5), the observable C is stable.

Proof that (3) implies (1). Let $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $(\sigma_l)_{l\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$ be sequences along which the ordinary iterated limits

$$L_1 = \lim_k \lim_l \varphi_C(\rho_k, \sigma_l), \qquad L_2 = \lim_l \lim_k \varphi_C(\rho_k, \sigma_l)$$

exist. Let $\mathcal{U}$ be a nonprincipal ultrafilter on $\mathbb{N}$, let $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$ and $\omega_2 = \lim_{l,\mathcal{U}} \sigma_l$, and put $\alpha_k = \lim_l \varphi_C(\rho_k, \sigma_l) = \omega_2(T_C\rho_k)$. By Lemma 2.4, the product $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$ equals L_2. Given $\varepsilon > 0$, choose n with $|\alpha_k - L_1| \leq \varepsilon$ for all $k \geq n$, and choose $\rho = \sum_{k\geq n} t_k\rho_k$ as in (3). Then we have

$$|L_2 - L_1| \leq |L_2 - \omega_2(T_C\rho)| + \Big|\sum_k t_k\alpha_k - L_1\Big| \leq 2\varepsilon.$$

Since $\varepsilon > 0$ is arbitrary, the two iterated limits agree, and the observable C is stable. □

In Example 6.2 below, we show that the order of idealization in (2) cannot be reversed and that the restriction to tails in (3) cannot be dropped: if Alice's system is idealized first, approximation by actual states can succeed for an unstable observable. For the order projection, by contrast, some idealized preparation has responses that no actual preparation reproduces to within any error less than 1, the largest possible error for a yes/no test, even on normal probes (see Example 6.3).

Example 6.2 (The order of idealization matters). Let $(e_k)_{k\geq 0}$ and $(f_l)_{l\geq 1}$ be orthonormal bases of H_1 and H_2, let Q_k be the projection onto the span of $f_1, \ldots, f_k$, and put

$$P = |e_0\rangle\langle e_0| \otimes I + \sum_{k\geq 1} |e_k\rangle\langle e_k| \otimes Q_k.$$

The operator P is a projection, and for $k \geq 1$ we have $\varphi_P(e_k, f_l) = 1$ if and only if $l \leq k$, so the projection P is unstable. On the other hand, for a density operator ρ on H_1 with diagonal entries $\rho_{kk} = \langle e_k|\rho|e_k\rangle$, the operator $T_P(\rho)$ is diagonal in the basis $(f_l)_{l\geq 1}$, with l-th diagonal entry $\rho_{00} + \sum_{k\geq l} \rho_{kk}$. These entries form a nonincreasing sequence in $[0, 1]$ whose first term is 1, and every such sequence $(c_l)_{l\geq 1}$ arises, from $\rho_{kk} = c_k - c_{k+1}$ and $\rho_{00} = \lim_l c_l$. The set of these diagonal operators is weak* closed, so for every state ω_1, the operator Z_{ω_1}, which lies in the weak* closure of $T_P(\mathfrak{D}(H_1))$, equals $T_P(\rho)$ for some actual ρ. Thus $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(P) = \omega_2(T_P\rho)$ exactly, for all states ω_2. For a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$ and $\omega_1 = \lim_{k,\mathcal{U}} e_k$, the simulating state is e_0; yet for $\omega_2 = \lim_{l,\mathcal{U}} f_l$ the two products differ: $(\omega_1 \otimes_{\mathsf{B}} \omega_2)(P) = 0$, whereas $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(P) = \omega_2(T_P e_0) = 1$.

Example 6.3 (Maximal failure for the order projection). Let $(e_k)_{k\in\mathbb{N}}$ and $(f_l)_{l\in\mathbb{N}}$ be an orthonormal basis of H_1 and an orthonormal sequence in H_2, respectively, let $P = P_{k<l}$ be the order projection, the projection onto the closed span of the vectors $e_k \otimes f_l$ with $k < l$, and let $\omega_1 = \lim_{k,\mathcal{U}} e_k$ for a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$. The effective observables $T_P(e_k) = \sum_{l>k} |f_l\rangle\langle f_l|$ decrease to 0 strongly, although each of them has norm 1; hence the operator Z_{ω_1} vanishes, and $(\omega_1 \otimes_{\mathsf{A}} \sigma)(P) = 0$ for every normal probe σ. For every $\rho \in \mathfrak{D}(H_1)$, however, with diagonal entries $\rho_{kk} = \langle e_k|\rho|e_k\rangle$, we have $T_P(\rho) =$

$\sum_k \rho_{kk} \sum_{l>k} |f_l\rangle\langle f_l|$, and $\mathrm{Tr}(f_N\, T_P \rho) = \sum_{k<N} \rho_{kk}$ tends to 1 as $N \to \infty$. Thus no actual preparation reproduces the responses of ω_1 to within any error less than 1, even on normal probes.

Remark 6.4 (Why degree one). In real-valued model theory, every type of a stable formula over a model, that is, every consistent assignment of a value of the formula to each parameter from the model, has a definition: a function on the parameters that gives these values and is a uniform limit of polynomials, or of lattice expressions, in instances of the formula [BYU10]. The instances are the functions obtained from the formula by fixing its first variable at elements of the model; for the formula $\varphi(x, y)$ interpreted by the expectation φ_C of an observable $C \in B(H_1 \otimes H_2)$, they are the actual responses $\varphi_C(\rho, \cdot\,)$, for $\rho \in \mathfrak{D}(H_1)$. Since φ_C is affine in the state, a convex combination of instances is again an instance: for density operators $\rho_1, \dots, \rho_m \in \mathfrak{D}(H_1)$ and weights $t_1, \dots, t_m \geq 0$ with $t_1 + \cdots + t_m = 1$, we have $\sum_{k=1}^m t_k \varphi_C(\rho_k, \cdot\,) = \varphi_C(\sum_{k=1}^m t_k \rho_k, \cdot\,)$. Thus in Theorem 6.1 the approximating "formula" is an instance $\varphi_C(\rho, \cdot\,)$ whose parameter ρ is a physical preparation, and polynomials of degree one suffice.

6.1 A dimension-free Monte Carlo bound

Mazur's theorem says nothing about the size of the simulating mixture in Theorem 6.1. In Theorem 6.5 below, we show that when the expectation of an observable factors through a Hilbert space, a sampling argument does: the mixture can be taken to be an empirical average of independent draws, and the law of large numbers in the Hilbert space bounds the error.

Theorem 6.5 (Monte Carlo simulation). *Let $C \in B(H_1 \otimes H_2)$ be given, let $\mathcal{H}$ be a Hilbert space, and let $a\colon \mathcal{T}(H_1) \to \mathcal{H}$ and $b\colon \mathcal{T}(H_2) \to \mathcal{H}$ be bounded maps on the spaces of trace-class operators, the first linear and the second conjugate-linear, such that the bilinear form m_C of C (see Definition 2.1) satisfies $m_C(\tau, \tau') = \langle a(\tau), b(\tau')\rangle$ for all $\tau \in \mathcal{T}(H_1)$ and $\tau' \in \mathcal{T}(H_2)$. Let $(\rho_k)_{k\in\mathbb{N}}$ be a sequence in $\mathfrak{D}(H_1)$, let $\mathcal{U}$ be a nonprincipal ultrafilter on $\mathbb{N}$, and let $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$. Then the observable C is stable, and for every $n \in \mathbb{N}$, every $\varepsilon > 0$ and every positive integer $M \geq (2R_a\|b\|/\varepsilon)^2$, where $R_a = \sup_{\rho\in\mathfrak{D}(H_1)} \|a(\rho)\|$ and $\|b\| = \sup_{\|\tau'\|_1\leq 1} \|b(\tau')\|$, there are indices $k_1, \dots, k_M \geq n$, not necessarily distinct, such that the equal-weight mixture $\bar\rho = \frac{1}{M}\sum_{j=1}^M \rho_{k_j}$ satisfies*

$$\big|(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) - \omega_2\big(T_C\bar\rho\big)\big| \leq \varepsilon \qquad \text{for all states } \omega_2 \text{ of } B(H_2).$$

Proof. For $\xi \in \mathcal{H}$, let $B_\xi \in B(H_2)$ be the operator with $\mathrm{Tr}(\tau' B_\xi) = \langle \xi, b(\tau')\rangle$ for all $\tau' \in \mathcal{T}(H_2)$. The map $\xi \mapsto B_\xi$ is linear, satisfies $\|B_\xi\| \leq \|b\|\,\|\xi\|$, and sends $a(\rho)$ to $T_C(\rho)$. Let ξ^* be the weak limit of $a(\rho_k)$ along $\mathcal{U}$. For a normal state σ, Lemma 2.4 yields $\mathrm{Tr}(\sigma Z_{\omega_1}) = \lim_{k,\mathcal{U}} \langle a(\rho_k), b(\sigma)\rangle = \langle \xi^*, b(\sigma)\rangle$, which identifies Z_{ω_1} with B_{ξ^*}. By Mazur's theorem, there is a finite convex combination $\xi = \sum_{k\geq n} t_k a(\rho_k)$ with $\|\xi^* - \xi\| \leq \varepsilon/(2\|b\|)$. Let $\kappa_1, \dots, \kappa_M$ be independent random indices with $\Pr[\kappa_j = k] = t_k$. Since the summands are independent with expectation ξ, we have

$$\mathbb{E}\Big\|\xi - \frac{1}{M}\sum_{j=1}^M a(\rho_{\kappa_j})\Big\|^2 = \frac{1}{M}\,\mathbb{E}\|a(\rho_{\kappa_1}) - \xi\|^2 \leq \frac{R_a^2}{M},$$

so some realization $k_1, \dots, k_M$ satisfies $\|\xi - \frac{1}{M}\sum_j a(\rho_{k_j})\| \leq R_a/\sqrt{M} \leq \varepsilon/(2\|b\|)$. Hence $\|Z_{\omega_1} - T_C\bar\rho\| = \|B_{\xi^* - a(\bar\rho)}\| \leq \varepsilon$. Stability follows from Proposition 4.4. □

Corollary 6.6. *In the setting of Theorem 6.5, for the swap operator F, which exchanges two systems with the same Hilbert space K, and for the equality tests $[A = B]$ between self-adjoint operators A and B with pure point spectra (see Section 4), some equal-weight mixture of $\lceil 4/\varepsilon^2 \rceil$ preparations ρ_k with $k \geq n$, not necessarily distinct, simulates the state ω_1 within ε, uniformly over all probes. For the swap test P_+ and the Hong–Ou–Mandel projection P_-, equal-weight mixtures of $\lceil 1/\varepsilon^2 \rceil$ such preparations suffice. These bounds do not depend on the dimension of the systems.*

Proof. For F, take for $\mathcal{H}$ the space $S_2(K)$ of Hilbert–Schmidt operators on K, with $a(\tau) = \tau$ and $b(\tau') = \tau'^*$, whence $\langle \tau, \tau'^* \rangle_{S_2} = \operatorname{Tr}(\tau'\tau)$; since the Hilbert–Schmidt norm satisfies $\|\tau\|_2 \leq \|\tau\|_1$, we have $R_a \leq 1$ and $\|b\| \leq 1$. On states, the slice maps satisfy $T_{P_\pm}(\rho) = \frac{1}{2}(I \pm T_F(\rho))$, and correspondingly, writing $Z^F_{\omega_1}$ and $Z^{P_\pm}_{\omega_1}$ for the operators Z_{ω_1} of F and of $P_\pm$, we have $Z^{P_\pm}_{\omega_1} = \frac{1}{2}(I \pm Z^F_{\omega_1})$, so the error for $P_\pm$ is half the error for F. For $[A = B]$, take $\mathcal{H} = \ell^2$, $a(\tau) = (\operatorname{Tr} \tau \Pi^A_\lambda)_\lambda$ and $b(\tau') = (\overline{\operatorname{Tr} \tau' \Pi^B_\lambda})_\lambda$, where Π^A_λ and Π^B_λ are the spectral projections of A and B for the eigenvalue λ; then $\|a(\tau)\|_2 \leq \sum_\lambda |\operatorname{Tr} \tau \Pi^A_\lambda| \leq \|\tau\|_1$, and similarly for b. □

If, in the setting of Theorem 6.5, the slice map T_C factors through a reflexive Banach space of Rademacher type $r > 1$ in place of the Hilbert space $\mathcal{H}$, the same argument, with the type-r form of Maurey's empirical method [Pis81], yields mixtures of size $M = O(\varepsilon^{-r/(r-1)})$ at accuracy ε. Here a Banach space has *Rademacher type* r if, for every positive integer m and all vectors $x_1, \ldots, x_m$ of the space, the average of $\|\pm x_1 \pm \cdots \pm x_m\|$ over the 2^m choices of signs is at most a constant multiple, independent of m, of $(\|x_1\|^r + \cdots + \|x_m\|^r)^{1/r}$. Reflexivity alone, which characterizes stability, yields no rate through this argument (see Question (Q5)).

6.2 Historical remarks

That the definitions of the types of a stable formula over a model, in real-valued model theory, are uniform limits of polynomials or lattice expressions in instances follows from the Stone–Weierstrass theorem (see [BYU10]); Ben Yaacov observed that, by Mazur's theorem, averages of instances already suffice [BY14]. Pták's lemma [Ptá63] is the classical source of convex means in the study of weak compactness. The sampling argument in the proof of Theorem 6.5 is Maurey's empirical method, published by Pisier [Pis81].

7 Probability: Fubini's theorem for finitely additive measures

Stability, as we show in this section, is a Fubini property: for any two probability measures, one on the set of preparations of each system, the two iterated integrals of the expectation against them agree. The property concerns *finitely additive* probability measures only, since for countably additive ones it always holds (see Proposition 7.2). We also relate the instability of the order projection, the yes/no test of whether, in two orthonormal sequences, the index of Bob's vector exceeds that of Alice's (see Theorem 4.7), to the noncommutativity of ultrafilter products.

7.1 Fubini's theorem for means

A *mean* on a set Γ is a finitely additive probability measure defined on all subsets of Γ; integration identifies such a measure with a positive linear functional on the bounded complex-valued functions on Γ that takes the value 1 at the constant function 1, and we

denote the measure and the functional by the same letter μ. We write $\int g\,d\mu$ for $\mu(g)$. The *barycenter* of a mean μ on $\mathfrak{D}(H)$ is the state $\omega_\mu(A) = \int \mathrm{Tr}(\rho A)\,d\mu(\rho)$ of $B(H)$.

Theorem 7.1 (Fubini's theorem for means). *An observable $C \in B(H_1 \otimes H_2)$, where H_1 and H_2 are the Hilbert spaces of Alice's and Bob's systems, is stable if and only if, for all means μ on $\mathfrak{D}(H_1)$ and ν on $\mathfrak{D}(H_2)$, its expectation $\varphi_C(\rho,\sigma) = \mathrm{Tr}\big((\rho\otimes\sigma)C\big)$ satisfies*

$$\int\Big(\int \varphi_C(\rho,\sigma)\,d\mu(\rho)\Big)d\nu(\sigma) = \int\Big(\int \varphi_C(\rho,\sigma)\,d\nu(\sigma)\Big)d\mu(\rho).$$

Proof. The idea is that, since the expectation is affine in each state, iterated integrals depend only on barycenters, whence they are the values at C of the two products of the barycenters (see Definition 2.3). The slice map S_C sends each preparation σ of Bob's system to the effective observable $S_C(\sigma)$ that it leaves on Alice's side, and the expectation $\varphi_C(\rho,\sigma) = \mathrm{Tr}(\rho\, S_C\sigma)$ is affine in ρ. Therefore the inner integral $\int \varphi_C(\rho,\sigma)\,d\mu(\rho)$ equals $\omega_\mu(S_C\sigma) = \mathrm{Tr}(\sigma Z_{\omega_\mu})$, where Z_{ω_μ} is the effective observable that the barycenter ω_μ leaves on Bob's side, and integration against ν yields $(\omega_\mu \otimes_{\mathsf{A}} \omega_\nu)(C)$, the product with Alice's system idealized first. Symmetrically, the other iterated integral is $(\omega_\mu \otimes_{\mathsf{B}} \omega_\nu)(C)$. Every state ω is a barycenter: by Lemma 2.2, there is a net (ρ_i) of normal states, indexed by a directed set, with limit ω, and if $\mathcal{W}$ is an ultrafilter on its index set containing all tails, then $\mu(g) = \lim_{i,\mathcal{W}} g(\rho_i)$ defines a mean with barycenter ω. Thus the identity holds for all means if and only if the two products agree on all pairs of states, which is condition (5) of Theorem 5.1, that is, if and only if C is stable. □

7.2 Countably additive measures

The means in Theorem 7.1 cannot be replaced by countably additive probability measures, because, as we show in Proposition 7.2 below, these never detect instability.

Proposition 7.2 (Countably additive measures never detect instability). *Let μ be a countably additive Borel probability measure on $\mathfrak{D}(H_1)$, with the trace norm, and let ν be a mean on $\mathfrak{D}(H_2)$. Then for every $C \in B(H_1 \otimes H_2)$, the two iterated integrals of φ_C against μ and ν agree.*

Proof. The barycenter of μ is the normal state given by the Bochner integral $\int \rho\,d\mu(\rho) \in \mathfrak{D}(H_1)$. By the proof of Theorem 7.1, the two iterated integrals are the values at C of the two products of the barycenters, and these values agree by Corollary 2.5. □

Instability is therefore a purely finitely additive phenomenon: it requires that both systems be idealized by genuinely finitely additive averages.

7.3 Ultrafilter products

For nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$, let $\mathcal{U}\otimes\mathcal{V}$ be the ultrafilter on $\mathbb{N}^2$ consisting of the sets $Q \subseteq \mathbb{N}^2$ with $\{k : \{l : (k,l) \in Q\} \in \mathcal{V}\} \in \mathcal{U}$. Let $C \in B(H_1 \otimes H_2)$ be given, let $(\rho_k)_{k\in\mathbb{N}}$ and $(\sigma_l)_{l\in\mathbb{N}}$ be sequences in $\mathfrak{D}(H_1)$ and $\mathfrak{D}(H_2)$, and let $\omega_1 = \lim_{k,\mathcal{U}} \rho_k$ and $\omega_2 = \lim_{l,\mathcal{V}} \sigma_l$ be their ultralimit states. By Lemma 2.4, the product $(\omega_1 \otimes_{\mathsf{B}} \omega_2)(C)$, with Bob's system idealized first, is the limit of $\varphi_C(\rho_k,\sigma_l)$ along $\mathcal{U}\otimes\mathcal{V}$, and the product $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$ is the limit along the transpose of $\mathcal{V}\otimes\mathcal{U}$, its image under the map $(l,k) \mapsto (k,l)$. The tensor square of a single nonprincipal ultrafilter $\mathcal{W}$ on $\mathbb{N}$ is not symmetric: the set $\{(k,l) : k < l\}$ belongs to $\mathcal{W}\otimes\mathcal{W}$ but not to the transpose of $\mathcal{W}\otimes\mathcal{W}$. The order projection of Theorem 4.7

is built from orthonormal sequences $(e_k)_{k\in\mathbb{N}}$ in H_1 and $(f_l)_{l\in\mathbb{N}}$ in H_2: it projects onto the closed span of the vectors $e_k \otimes f_l$ with $k < l$, whence its expectation on the pairs (e_k, f_l) is the indicator of the set $\{(k,l) : k < l\}$. For $\omega_1 = \lim_{k,\mathcal{W}} e_k$ and $\omega_2 = \lim_{l,\mathcal{W}} f_l$, the product $\otimes_{\mathsf{B}}$, with Bob's system idealized first, therefore takes the value 1 at the order projection, whereas the product $\otimes_{\mathsf{A}}$ takes the value 0. The ultrafilter $\mathcal{W}$ corresponds to a $\{0,1\}$-valued mean on $\mathbb{N}$, which gives measure 1 exactly to the sets in $\mathcal{W}$; by Proposition 7.2, no countably additive version of this example exists.

7.4 Means, mixtures and samples

Finitely supported means on $\mathfrak{D}(H_1)$ correspond to the finite mixtures a laboratory can implement. In this language, Theorem 6.1 says that, for a stable observable and every $\varepsilon > 0$, a mean μ on $\mathfrak{D}(H_1)$ in Theorem 7.1 can be replaced by a finitely supported mean, with an error at most ε uniformly in the mean ν on $\mathfrak{D}(H_2)$, because the iterated integrals depend on μ only through its barycenter (see the proof of Theorem 7.1). For observables whose expectation factors through a Hilbert space, Theorem 6.5 replaces Mazur's existence statement by an empirical average of independent draws.

7.5 Historical remarks

Pták's combinatorial lemma [Ptá63] derives weak compactness from the existence of convex means that interchange limits, and the Fubini form for finitely additive means was studied by Simons [Sim68] and by Levy [Lev07]. In model theory, finitely additive probability measures on definable sets, the Keisler measures [Kei87], admit two products according to the order of integration; the commutation of these products is tied to stability and to generic stability, a property of a single type or measure in a theory that need not be stable; in theories without the independence property, the NIP theories, a global type or measure, one over a very large saturated model, is generically stable exactly when it is both definable over and finitely satisfiable in some model of smaller cardinality [HPP08, HP11]. The two products $\otimes_{\mathsf{A}}$ and $\otimes_{\mathsf{B}}$ of Section 2, with Alice's and with Bob's system idealized first, respectively, are the operator-algebraic counterparts of the two products of Keisler measures (see the proof of Theorem 7.1).

8 Independence: Harrington's lemma and the independent product

On stable observables, two idealized preparations, one of each system, combine independently in a unique and symmetric way, while on an unstable observable the two orders of idealization disagree for some pair of idealized preparations; we prove both statements in Theorem 8.2 below. The model-theoretic counterpart of these statements is the theory of independence. In a stable theory, independence over a model is symmetric, and the key to symmetry is Harrington's lemma: the definitions of two types over a model commute [HH84]. We recall the origins of abstract independence briefly, explain that Harrington's lemma asserts that two iterated ultrafilter limits can be exchanged, and then return to observables.

8.1 Abstract independence

The axiomatic study of independence is older than model theory: it begins with Steinitz's exchange argument for transcendence bases and with van der Waerden's abstract dependence relations, and leads to matroids and pregeometries, closure operators with the ex-

change property of linear span (see the historical remarks at the end of this section). In model theory, Shelah's forking [She90] carries independence beyond pregeometries. Its negation, nonforking, singles out the extensions of a type to larger sets of parameters whose realizations are independent of the new parameters; nonforking independence is not given by a closure operator, yet in a stable theory it is symmetric and has existence and uniqueness of independent extensions over models. Lascar and Poizat [LP79] described it through definability of types and through heirs and coheirs of types over a model: the extensions in which every formula, once its new parameters are replaced by suitable parameters from the model, belongs to the original type, and those in which every formula is satisfied in the model, respectively. Harnik and Harrington [HH84] characterized stable theories as those carrying an abstract independence relation with these properties.

8.2 Harrington's lemma as an exchange of limits

We recall the two-valued setting of Iovino's notes on stability theory [Iov, Chapter 4]. Let $\mathcal{M}$ be a model of a stable complete first-order theory, let p and q be complete types over $\mathcal{M}$ in finite tuples of variables x and y, and let $\varphi(x,y)$ be a formula with parameters from $\mathcal{M}$. Stability provides definitions $d_p\varphi(y)$ and $d_q\varphi(x)$ over $\mathcal{M}$: for every tuple b from $\mathcal{M}$, the formula $\varphi(x,b)$ belongs to p if and only if $d_p\varphi(b)$ holds in $\mathcal{M}$, and symmetrically for q. In Section 8.3 below, a type over $\mathcal{M}$ will correspond to an idealized state, and realizing one coheir before another to idealizing the corresponding system first, that is, to taking the limit over its states as the inner limit.

Fact 8.1 (Harrington's lemma [HH84]). *Let $\mathcal{M}$ be a model of a stable complete first-order theory, let p and q be complete types over $\mathcal{M}$ in finite tuples of variables x and y, and let $\varphi(x,y)$ be a formula with parameters from $\mathcal{M}$. Then the definitions $d_p\varphi(y)$ and $d_q\varphi(x)$ of p and q over $\mathcal{M}$ satisfy $d_p\varphi(y) \in q$ if and only if $d_q\varphi(x) \in p$.*

Harrington's lemma is the key step in the proof that nonforking independence is symmetric in stable theories, and hence in the Harnik–Harrington characterization [HH84]. For us, the point is that Harrington's lemma asserts that two iterated limits can be exchanged: every type over $\mathcal{M}$ is finitely satisfiable in $\mathcal{M}$, so the type p is the trace of an ultrafilter $\mathcal{U}$ on the tuples from $\mathcal{M}$ of the length of x, in the sense that a formula $\theta(x)$ with parameters from $\mathcal{M}$ belongs to p if and only if the set of its solutions in $\mathcal{M}$ belongs to $\mathcal{U}$. Let $\mathcal{V}$ be an ultrafilter with trace q, and regard φ as a $\{0,1\}$-valued function on pairs of tuples from $\mathcal{M}$. Then the two sides of Harrington's lemma take the form

$$d_p\varphi(y) \in q \iff \lim_{b,\mathcal{V}} \lim_{a,\mathcal{U}} \varphi(a,b) = 1, \qquad d_q\varphi(x) \in p \iff \lim_{a,\mathcal{U}} \lim_{b,\mathcal{V}} \varphi(a,b) = 1,$$

so, read in the language of types, the agreement of the two iterated ultrafilter limits of Corollary 3.3 is exactly what Harrington's lemma asserts. The computation takes place entirely in $\mathcal{M}$ and involves no heirs, no saturated models, those that realize all types over sets of parameters of smaller cardinality than the model, and no realizations. The same computation explains the two orders of realization: let b realize q, and let a realize the coheir of p over $\mathcal{M}b$, the set $\mathcal{M}$ with the entries of b adjoined, determined by $\mathcal{U}$, that is, the type over $\mathcal{M}b$ that contains a formula exactly when the set of its solutions in $\mathcal{M}$ belongs to $\mathcal{U}$. Then the formula $\varphi(a,b)$ holds if and only if $\lim_{a',\mathcal{U}} \lim_{b',\mathcal{V}} \varphi(a',b') = 1$, and realizing p first and then the coheir of q over $\mathcal{M}a$ determined by $\mathcal{V}$ exchanges the limits. Thus the two orders of realizing coheirs yield the same type exactly when the two iterated limits agree, and the symmetry of independence reduces to the agreement of iterated ultrafilter limits. The proofs of these statements, and of the Harnik–Harrington characterization, are given

in the same notes [Iov, Chapter 4]; the real-valued versions, for formulas with values in a compact interval, are proved in Iovino's thesis [Iov94] (see also [BYU10]).

8.3 The independent product of idealized preparations

The correspondence with Section 2 is as follows. A type over a model corresponds to an idealized preparation, and the trace of an ultrafilter to an ultralimit state. Realizing Bob's coheir first and Alice's over it computes the inner limit over Bob's states, just as the product $\otimes_{\mathsf{B}}$ of Definition 2.3, in which Bob's system is idealized first, computes it; the other order computes $\otimes_{\mathsf{A}}$. Under this correspondence, the conclusion of Harrington's lemma becomes condition (5) of Theorem 5.1, the agreement of the two products on every pair of states. The analogy is one of form: the theorem below is proved directly from Theorem 5.1, without model theory. To state the resulting symmetry, let $\mathcal{O}$ be the set of stable observables in $B(H_1 \otimes H_2)$, let $F_{12}\colon H_1 \otimes H_2 \to H_2 \otimes H_1$ be the unitary with $F_{12}(\psi \otimes \chi) = \chi \otimes \psi$ for $\psi \in H_1$ and $\chi \in H_2$, and for states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, let $\omega_1 \otimes \omega_2$ be the functional on $\mathcal{O}$ that takes at each $C \in \mathcal{O}$ the common value of the two products, which agree at C by Theorem 5.1:

$$(\omega_1 \otimes \omega_2)(C) = (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = (\omega_1 \otimes_{\mathsf{B}} \omega_2)(C) \qquad (C \in \mathcal{O}).$$

In Proposition 10.1 below, we show that the set $\mathcal{O}$ is a norm-closed self-adjoint subspace of $B(H_1 \otimes H_2)$ containing I.

Theorem 8.2 (The independent product). *Let $\mathcal{O}$ be the set of stable observables in $B(H_1 \otimes H_2)$, and for states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, let $\omega_1 \otimes \omega_2$ be their independent product, the functional on $\mathcal{O}$ that takes at each $C \in \mathcal{O}$ the common value of the two products of Definition 2.3. Then the following hold:*

(a) (Existence.) *For all states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, the functional $\omega_1 \otimes \omega_2$ is a state of $\mathcal{O}$: it is linear, positive on positive elements of $\mathcal{O}$, and takes the value 1 at I.*

(b) (Extension.) *If states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$ are given by density operators ρ and σ, then $(\omega_1 \otimes \omega_2)(C) = \operatorname{Tr}((\rho \otimes \sigma)C)$ for all $C \in \mathcal{O}$.*

(c) (Uniqueness.) *For every $C \in \mathcal{O}$, the assignment $(\omega_1, \omega_2) \mapsto (\omega_1 \otimes \omega_2)(C)$ is separately weak* continuous, and it is the only separately weak* continuous function on $\mathcal{S}(B(H_1)) \times \mathcal{S}(B(H_2))$ that takes the value $\operatorname{Tr}((\rho \otimes \sigma)C)$ at every pair of normal states given by density operators ρ and σ.*

(d) (Symmetry.) *Let $F_{12}\colon H_1 \otimes H_2 \to H_2 \otimes H_1$ be the unitary that exchanges the two factors. For every $C \in \mathcal{O}$, the observable $F_{12} C F_{12}^*$ on $H_2 \otimes H_1$ is stable, and $(\omega_2 \otimes \omega_1)(F_{12} C F_{12}^*) = (\omega_1 \otimes \omega_2)(C)$ for all states ω_1 of $B(H_1)$ and ω_2 of $B(H_2)$, the functional $\omega_2 \otimes \omega_1$ being defined in the same way on the stable observables on $H_2 \otimes H_1$.*

Moreover, an observable $C \in B(H_1 \otimes H_2)$ is stable if and only if the two products $\otimes_{\mathsf{A}}$ and $\otimes_{\mathsf{B}}$ agree at C for all pairs of states.

Proof. (a) The slice map S_C, which sends each preparation of Bob's system to the effective observable that it leaves on Alice's side, is linear in C, and so is $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \omega_2(Z_{\omega_1})$, where $Z_{\omega_1} = S_C^*(\omega_1)$ is the effective observable that ω_1 leaves on Bob's side. If $C \geq 0$, then

the operator $S_C(\sigma)$ is positive for every positive $\sigma \in \mathfrak{T}(H_2)$, so $\mathrm{Tr}(\sigma Z_{\omega_1}) = \omega_1(S_C\sigma) \geq 0$ for every such σ; hence the operator Z_{ω_1} is positive, and $(\omega_1 \otimes_{\mathsf{A}} \omega_2)(C) = \omega_2(Z_{\omega_1}) \geq 0$. For $C = I$, we have $S_I(\sigma) = \mathrm{Tr}(\sigma)I$ and $Z_{\omega_1} = I$.

(b) This is Corollary 2.5.

(c) Both statements come from the step of the proof of Theorem 5.1 in which condition (5), the agreement of the two products, is shown to be equivalent to condition (6), the existence of a separately continuous extension.

(d) Write $C' = F_{12}CF_{12}^*$. For density operators we have $\varphi_{C'}(\sigma, \rho) = \varphi_C(\rho, \sigma)$, so the observable C' is stable by Lemma 3.4(c). Let (ρ_i) and (σ_j) be nets of normal states of $B(H_1)$ and of $B(H_2)$, indexed by directed sets, with limits ω_1 and ω_2. By Lemma 2.4, the product $(\omega_2 \otimes_{\mathsf{B}} \omega_1)(C')$, in which the second factor ω_1 is idealized first, equals $\lim_j \lim_i \varphi_{C'}(\sigma_j, \rho_i) = \lim_j \lim_i \varphi_C(\rho_i, \sigma_j) = (\omega_1 \otimes_{\mathsf{A}} \omega_2)(C)$. Since C' and C are stable, the two sides are the independent products $(\omega_2 \otimes \omega_1)(C')$ and $(\omega_1 \otimes \omega_2)(C)$, and (d) follows.

The final statement is the equivalence, in Theorem 5.1, of condition (1), stability, and condition (5), the agreement of the two products. □

In the language of independence relations, (a) is the existence of an independent combination, (c) is its uniqueness (stationarity), and (d) is its symmetry. Outside $\mathcal{O}$ the independent combination is lost: by Theorem 5.1, an unstable observable admits no separately continuous extension of its expectation to idealized states, and the two orders of idealization disagree on some pair of idealized states. For the order projection, whose expectation on pairs of vectors from two orthonormal sequences records whether Alice's index is smaller than Bob's, the two orders of idealization give 1 with Bob's system idealized first and 0 with Alice's (see Theorem 4.7).

8.4 Historical remarks

Steinitz [Ste10] introduced transcendence bases of field extensions and proved, by the exchange argument that also gives the dimension of a vector space, that any two of them have the same cardinality. Van der Waerden [vdW30] observed that linear and algebraic dependence obey the same few rules, and derived the existence of bases and the invariance of their cardinality from an abstract dependence relation whose central axiom is exchange. Whitney [Whi35] isolated these rules as *matroids*, with equivalent axiomatizations by independent sets, bases, circuits and rank; Nakasawa arrived at the theory independently at about the same time [NK09], and Birkhoff [Bir35] and Mac Lane [ML36] related it to lattices and to projective geometry. Crapo and Rota [CR70] developed matroids as *combinatorial geometries*; in their terminology a closure operator with the exchange property is a *pregeometry*, the term still used in model theory. Baldwin and Lachlan [BL71] showed that algebraic closure in a strongly minimal set is a pregeometry, and that every model of a strongly minimal theory therefore has a dimension. Here a set is strongly minimal if it is an infinite definable set whose definable subsets are finite or cofinite in every elementary extension, a theory is strongly minimal if the universe of its models is, and the algebraic closure of a set consists of the elements that satisfy a formula with parameters from the set and with only finitely many solutions. Kim and Pillay [KP97] extended the Harnik–Harrington characterization to simple theories, a class that properly contains the stable theories.

9 Combinatorics: polynomial versus exponential growth

This section is devoted to Theorem B of the Introduction, which settles the problem of discrimination: it bounds the number of preparations that a given number of adaptive probes can tell apart, and it characterizes the observables whose responses can be learned with a bounded number of mistakes. We show that unstable observables encode binary decision trees of every finite height, that bounded tree height is a uniform form of stability characterized through ultrapowers, and that it is equivalent both to polynomial growth of the number of preparations resolved and to a finite bound, uniform in the threshold, on the mistakes made in a learning game with a fixed threshold and margin.

Throughout the section, we fix a self-adjoint observable C on $H_1 \otimes H_2$, where H_1 and H_2 are the Hilbert spaces of Alice's and Bob's systems; its expectation $\varphi_C(\rho,\sigma) = \mathrm{Tr}\big((\rho\otimes\sigma)C\big)$, for density operators $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$, is therefore real. For a general observable, the results apply to its hermitian parts (see Definition 9.5). For every integer $n \geq 0$, we write 2^n for the set of 0–1 sequences of length n, and we put $2^{<n} = \bigcup_{m<n} 2^m$; the sets of infinite and of finite 0–1 sequences are written $2^{\mathbb{N}}$ and $2^{<\mathbb{N}}$, respectively. For a finite 0–1 sequence s and a finite or infinite 0–1 sequence η, we write $s \sqsubseteq \eta$ if s is an initial segment of η, and $s^\frown\iota$ for s followed by $\iota \in \{0,1\}$.

9.1 From ladders to decision trees

Definition 9.1 (Ladders and trees). Let $\delta > 0$ be given.

(a) For a positive integer N, a *δ-ladder of length N* for C consists of states $\rho_1,\dots,\rho_N \in \mathfrak{D}(H_1)$ and $\sigma_1,\dots,\sigma_N \in \mathfrak{D}(H_2)$, a real number γ and a sign $\theta \in \{\pm 1\}$ such that

$$\theta\big(\varphi_C(\rho_i,\sigma_j) - \gamma\big) \geq \delta/2 \ \text{ if } i \leq j, \qquad \theta\big(\varphi_C(\rho_i,\sigma_j) - \gamma\big) \leq -\delta/2 \ \text{ if } i > j.$$

Infinite ladders are defined in the same way.

(b) For an integer $n \geq 0$, a *δ-tree of height n* for C consists of states $\rho_\eta \in \mathfrak{D}(H_1)$ for $\eta \in 2^n$ (the *leaves*), states $\sigma_s \in \mathfrak{D}(H_2)$ for $s \in 2^{<n}$ (the *nodes*), a real number γ and a sign $\theta \in \{\pm 1\}$ such that, whenever $s^\frown\iota \sqsubseteq \eta$, we have

$$\theta\big(\varphi_C(\rho_\eta,\sigma_s) - \gamma\big) \geq \delta/2 \ \text{ if } \iota = 1, \qquad \theta\big(\varphi_C(\rho_\eta,\sigma_s) - \gamma\big) \leq -\delta/2 \ \text{ if } \iota = 0.$$

A tree describes an adaptive measurement protocol. Bob starts at the root, the empty sequence $\emptyset$, and, at each node s, probes with σ_s; the expectation of C falls on one side of the threshold γ or the other, and Bob moves to $s^\frown 1$ or $s^\frown 0$ accordingly. If Alice prepared one of the leaves, then after n probes he knows which of the 2^n preparations ρ_η she made. Two leaves that first separate at the node s have responses to σ_s that differ by at least δ.

Proposition 9.2 (Infinite ladders from instability). *If sequences $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $(\sigma_l)_{l\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$ are such that the iterated limits $L_1 = \lim_k \lim_l \varphi_C(\rho_k,\sigma_l)$ and $L_2 = \lim_l \lim_k \varphi_C(\rho_k,\sigma_l)$ exist and differ, then for every δ with $0 < \delta < |L_1 - L_2|$ there is an infinite δ-ladder for C whose states are taken from these sequences.*

Proof. For any nonprincipal ultrafilters $\mathcal{U}$ and $\mathcal{V}$ on $\mathbb{N}$, the iterated limits along $\mathcal{U}$ and $\mathcal{V}$ equal the ordinary iterated limits L_1 and L_2. Apply Lemma 3.2 with $\varepsilon = (|L_1 - L_2| - \delta)/2$, and take $\gamma = (L_1 + L_2)/2$ and θ the sign of $L_1 - L_2$. □

The passage from ladders to trees rests on a single arithmetic observation, which we prove in Proposition 9.3 below: if each finite binary sequence is placed at the point of $(0,1)$ whose binary expansion consists of its digits complemented and followed by a final 1, then each leaf lies to the left of every node at which its branch takes the label 1, and to the right of every node at which it takes the label 0 (see Figure 9.1).

Proposition 9.3 (Ladders give trees). *For every $\delta > 0$ and every integer $n \geq 0$, every δ-ladder of length $2^{n+1}-1$ for C yields a δ-tree of height n for C, with the same threshold γ and sign θ, and with leaves and nodes taken from the ladder.*

Proof. For a 0–1 sequence s of length $|s| \leq n$, with digits $s(0), \dots, s(|s|-1)$, put

$$\mathrm{pos}(s) = \sum_{k<|s|} (1-s(k))2^{-(k+1)} + 2^{-(|s|+1)} \in (0,1).$$

The binary expansion of $\mathrm{pos}(s)$ is $1-s(0), \dots, 1-s(|s|-1), 1$, so the map pos is injective, and its image has $2^{n+1}-1$ elements. We claim that if $s^\frown\iota \sqsubseteq \eta$ with $\eta \in 2^n$, then the leaf η lies to the left of s, that is, $\mathrm{pos}(\eta) < \mathrm{pos}(s)$, if and only if $\iota = 1$. Indeed, with $k = |s|$, we have

$$\mathrm{pos}(\eta) - \mathrm{pos}(s) = -\iota\, 2^{-(k+1)} + \Delta, \qquad \Delta = \sum_{k<j<n} (1-\eta(j))2^{-(j+1)} + 2^{-(n+1)},$$

and $0 < \Delta \leq (2^{-(k+1)} - 2^{-n}) + 2^{-(n+1)} < 2^{-(k+1)}$. The bound $0 < \Delta < 2^{-(k+1)}$ is the point of the construction: the digits after position k can move $\mathrm{pos}(\eta)$ by less than the digit at position k, so the turn taken at s alone decides on which side of $\mathrm{pos}(s)$ the leaf lies.

Let $\rho_1, \dots, \rho_N$ and $\sigma_1, \dots, \sigma_N$, with $N = 2^{n+1}-1$, be the states of the ladder, and let π be the increasing bijection from the image of pos onto $\{1, \dots, N\}$. Put $\rho_\eta = \rho_{\pi(\mathrm{pos}(\eta))}$ for $\eta \in 2^n$ and $\sigma_s = \sigma_{\pi(\mathrm{pos}(s))}$ for $s \in 2^{<n}$. If $s^\frown\iota \sqsubseteq \eta$ and $\iota = 1$, then $\pi(\mathrm{pos}(\eta)) < \pi(\mathrm{pos}(s))$, and the ladder yields $\theta(\varphi_C(\rho_\eta, \sigma_s) - \gamma) \geq \delta/2$; if $\iota = 0$, then $\pi(\mathrm{pos}(\eta)) > \pi(\mathrm{pos}(s))$, and the opposite inequality follows. □

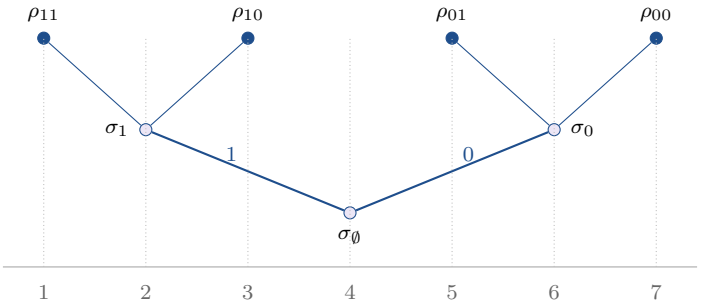


Figure 9.1: The map pos for $n = 2$. The seven sequences of length at most 2 are placed at the positions $1, \dots, 7$ of a ladder of length 7, ordered by pos; the tree grows upward. Each leaf lies to the left of the nodes at which its branch takes the label 1, so the ladder inequalities for $i \leq j$ are the tree inequalities for $\iota = 1$.

Theorem 9.4 (Unstable observables encode trees). *If C is unstable, then there is $\delta > 0$ such that for every $n \in \mathbb{N}$, there is a δ-tree of height n for C. More precisely, if sequences $(\rho_k)_{k\in\mathbb{N}}$ in $\mathfrak{D}(H_1)$ and $(\sigma_l)_{l\in\mathbb{N}}$ in $\mathfrak{D}(H_2)$ are such that the two iterated limits of $\varphi_C(\rho_k, \sigma_l)$ exist and differ by $\Delta > 0$, then for every δ with $0 < \delta < \Delta$ and every $n \in \mathbb{N}$, there is a δ-tree of height n for C whose leaves and nodes are terms of these sequences; in particular, they are pure states when the terms of the sequences are pure.*

Proof. Combine Propositions 9.2 and 9.3. □

For the order projection, built from two orthonormal sequences $(e_k)_{k\in\mathbb{N}}$ and $(f_l)_{l\in\mathbb{N}}$, and for the comparison $[A < B]$ of two self-adjoint operators A and B with pure point spectra that are unbounded above, the values of the expectation on the vectors e_k and f_l and on eigenvectors of A and B, respectively, form 1-ladders of every length (see Theorem 4.7 and Proposition 4.6(3)); hence, by Proposition 9.3, there are 1-trees of every height whose leaves and nodes are the pure states given by these vectors. The converse of Theorem 9.4 fails (see Example 9.8 below): trees of unbounded height detect a uniform form of instability, which we now describe through ultrapowers.

9.2 Ultrapowers and uniform stability

Fix a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$. A sequence $x = (x_k)_{k\in\mathbb{N}}$ of states of Alice's system represents an element of the ultrapower of her preparations; it should be thought of as a family of experiments indexed by k, in which both parties vary their preparations together. The *ultrapower expectation* of C is the function

$$\Phi_C^{\mathcal{U}}(x, y) = \lim_{k,\mathcal{U}} \varphi_C(x_k, y_k) \qquad (x \in \mathfrak{D}(H_1)^{\mathbb{N}},\ y \in \mathfrak{D}(H_2)^{\mathbb{N}}).$$

It is a joint limit, not an iterated one, and so it differs in kind from the two products of idealized states in Section 2; its value at (x, y) is that, at the classes of x and y, of the bilinear form induced by $m_C(\tau, \tau') = \mathrm{Tr}((\tau \otimes \tau')C)$ on the Banach-space ultrapowers of the spaces $\mathcal{T}(H_1)$ and $\mathcal{T}(H_2)$ of trace-class operators [Hei80]. Ladders and trees for $\Phi_C^{\mathcal{U}}$ are defined as in Definition 9.1, with sequences in $\mathfrak{D}(H_1)^{\mathbb{N}}$ and $\mathfrak{D}(H_2)^{\mathbb{N}}$ in place of states. For $\delta > 0$, an *infinite δ-tree* for $\Phi_C^{\mathcal{U}}$ consists of leaves $x_\eta \in \mathfrak{D}(H_1)^{\mathbb{N}}$ for every infinite branch $\eta \in 2^{\mathbb{N}}$, nodes $y_s \in \mathfrak{D}(H_2)^{\mathbb{N}}$ for $s \in 2^{<\mathbb{N}}$, a real number γ and a sign θ, satisfying the inequalities of Definition 9.1(b), with $\Phi_C^{\mathcal{U}}$ in place of φ_C, whenever $s^\frown\iota \sqsubseteq \eta$.

Definition 9.5. A self-adjoint observable C is *uniformly stable* if for every $\delta > 0$, the lengths of δ-ladders for C are bounded. An arbitrary bounded observable C is *uniformly stable* if its hermitian parts $\frac{1}{2}(C + C^*)$ and $\frac{1}{2i}(C - C^*)$ are uniformly stable.

Theorem 9.6 (Uniform stability). *For every self-adjoint $C \in B(H_1 \otimes H_2)$, the following are equivalent:*

(1) *The observable C is uniformly stable.*

(2) *For every $\delta > 0$, the heights of δ-trees for C are bounded.*

(3) *The ultrapower expectation $\Phi_C^{\mathcal{U}}$ is stable for every nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$.*

(4) *The ultrapower expectation $\Phi_C^{\mathcal{U}}$ is stable for some nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$.*

If these conditions fail, then there is $\delta > 0$ such that for every nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, there is an infinite δ-tree for $\Phi_C^{\mathcal{U}}$. Every uniformly stable observable is stable.

Proof. The implication from (2) to (1) is Proposition 9.3.

Proof that (1) implies (3). Suppose that $\Phi_C^{\mathcal{U}}$ is unstable. The proof of Proposition 9.2 uses only Lemma 3.2, which holds for any bounded function on a product of sets, so it provides, for some $\delta > 0$, an infinite δ-ladder for $\Phi_C^{\mathcal{U}}$, consisting of sequences $x_i \in \mathfrak{D}(H_1)^{\mathbb{N}}$ and $y_j \in \mathfrak{D}(H_2)^{\mathbb{N}}$ for $i, j \in \mathbb{N}$, a threshold γ and a sign θ; we write $x_{i,k}$ and $y_{j,k}$ for the k-th entries of x_i and y_j. Fix a positive integer N. For each of the N^2 pairs $i, j \le N$, the set of k with $\theta(\varphi_C(x_{i,k}, y_{j,k}) - \gamma) \ge \delta/4$ (if $i \le j$), respectively $\le -\delta/4$ (if $i > j$), belongs to $\mathcal{U}$. For k in the intersection of these sets, the states $x_{i,k}$ and $y_{j,k}$ form a $\delta/2$-ladder of length N for C. Hence the observable C is not uniformly stable.

The implication from (3) to (4) is trivial.

Proof that (4) implies (2). The idea is to glue trees of increasing height into an infinite tree in the ultrapower, and to observe that an infinite tree has $2^{\aleph_0}$ separated leaves, which no separable dual can accommodate. Suppose that for some $\delta > 0$ there are δ-trees $(\rho_\eta^{(n)}, \sigma_s^{(n)}, \gamma_n, \theta_n)$ of every height n, and let $\mathcal{U}$ be any nonprincipal ultrafilter on $\mathbb{N}$. Pass to a set in $\mathcal{U}$ on which $\theta_n = \theta$ is constant, and let $\gamma = \lim_{n,\mathcal{U}} \gamma_n$, which exists because, for $n \ge 1$, the root of the n-th tree has leaves on both sides of γ_n, whence $|\gamma_n| \le \|C\|$. For an infinite branch $\eta \in 2^{\mathbb{N}}$, put $x_\eta = (\rho_{\eta|n}^{(n)})_n$, with $\eta|n$ the initial segment of length n, and for $s \in 2^{<\mathbb{N}}$, put $y_s = (\sigma_s^{(n)})_n$, with arbitrary entries for $n \le |s|$. If $s^\frown\iota \sqsubseteq \eta$, then the same holds with $\eta|n$ in place of η for all $n > |s|$, and passing to the limit, we obtain $\theta(\Phi_C^{\mathcal{U}}(x_\eta, y_s) - \gamma) \ge \delta/2$ if $\iota = 1$ and $\le -\delta/2$ if $\iota = 0$. This is an infinite δ-tree for $\Phi_C^{\mathcal{U}}$, which proves the statement of Theorem 9.6 on infinite δ-trees.

Now take for $\mathcal{U}$ the ultrafilter of (4), for which $\Phi_C^{\mathcal{U}}$ is stable. After an affine rescaling to values in $[0, 1]$, which multiplies gaps by $c = 1/(2\|C\|)$, the theorem of Iovino [Iov99b, Theorem 3.3] yields a reflexive Banach space E and maps $u\colon \mathfrak{D}(H_1)^{\mathbb{N}} \to E$ and $v\colon \mathfrak{D}(H_2)^{\mathbb{N}} \to E^*$, with values in the unit balls, such that $c(\Phi_C^{\mathcal{U}} + \|C\|) = \langle u(\cdot), v(\cdot)\rangle$. Let Y be the closed span in E^* of the countably many vectors $v(y_s)$. The space Y is separable and reflexive, so its dual Y^* is separable. If $\eta \ne \eta'$ first differ after the segment s, then $|\langle u(x_\eta) - u(x_{\eta'}), v(y_s)\rangle| \ge c\delta$, so the restrictions of the vectors $u(x_\eta)$ to Y are $2^{\aleph_0}$ points of Y^* at mutual distance at least $c\delta$, which is impossible in a separable space. This contradiction shows that the heights of δ-trees for C are bounded for every $\delta > 0$, which is (2).

Finally, an unstable observable has infinite ladders by Proposition 9.2, so it is not uniformly stable. □

Conditions (1), (2) and (3) of Theorem 9.6, but not condition (4), remain equivalent for

a bounded family $(C_k)_{k\in\mathbb{N}}$ of self-adjoint observables C_k on finite-dimensional spaces $H_1^{(k)} \otimes H_2^{(k)}$, with $\Phi^{\mathcal{U}}(x,y) = \lim_{k,\mathcal{U}} \operatorname{Tr}((x_k \otimes y_k)C_k)$ for sequences $x = (x_k)_{k\in\mathbb{N}}$ and $y = (y_k)_{k\in\mathbb{N}}$ with $x_k \in \mathfrak{D}(H_1^{(k)})$ and $y_k \in \mathfrak{D}(H_2^{(k)})$, when the ladders and trees of the family are those of its members and their lengths and heights are required to be bounded independently of k (if $C_k = 0$ for every odd k, then $\Phi^{\mathcal{U}} = 0$ for every $\mathcal{U}$ that contains the set of odd numbers, whatever the members of even index); this is the relevant notion for registers of growing size (see Section 11).

Proposition 9.7. *Every self-adjoint C in the minimal tensor product $B(H_1) \otimes_{\min} B(H_2)$, the norm closure of the finite sums of products $A \otimes B$ of bounded operators on H_1 and on H_2, is uniformly stable, and so is every self-adjoint $C \in B(H_1 \otimes H_2)$ for which there are a Hilbert space $\mathcal{H}$ and bounded maps $u\colon \mathfrak{D}(H_1) \to \mathcal{H}$ and $v\colon \mathfrak{D}(H_2) \to \mathcal{H}$ with $\varphi_C(\rho,\sigma) = \langle u(\rho), v(\sigma)\rangle$ for all $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$. In particular, the swap operator F, which exchanges two systems with the same Hilbert space, the projections $P_\pm$ of the swap test and of Hong–Ou–Mandel interference, and the equality tests $[A = B]$ of Section 4 are uniformly stable.*

Proof. For a finite sum $C' = \sum_{j=1}^{r} A_j \otimes B_j$ of products, with $A_j \in B(H_1)$ and $B_j \in B(H_2)$, the function $\Phi^{\mathcal{U}}_{C'}(x,y) = \sum_j \lim_{k,\mathcal{U}} \operatorname{Tr}(x_k A_j) \cdot \lim_{k,\mathcal{U}} \operatorname{Tr}(y_k B_j)$ is a separated sum, and the argument of Theorem 4.1 then shows that the function $\Phi^{\mathcal{U}}_C$ is stable. In the Hilbert case, we have $\Phi^{\mathcal{U}}_C(x,y) = \langle (u(x_k))_{\mathcal{U}}, (v(y_k))_{\mathcal{U}}\rangle$ in the ultrapower Hilbert space $\mathcal{H}^{\mathcal{U}}$, where $(u(x_k))_{\mathcal{U}}$ denotes the element of $\mathcal{H}^{\mathcal{U}}$ represented by the sequence $(u(x_k))_{k\in\mathbb{N}}$, and Proposition 4.4 applies. Conclude with Theorem 9.6. □

Example 9.8 (Stable but not uniformly stable). Let both H_1 and H_2 be the Hilbert space $\bigoplus_{n\geq 1} \mathbb{C}^n$, with orthonormal bases $(e_i^n)_{i\leq n}$ of the summands, which we call sectors; for instance, for $j \in \{0, \frac{1}{2}, 1, \frac{3}{2}, \dots\}$, the summand $\mathbb{C}^{2j+1}$ carries the spin-j representation, the irreducible representation of dimension $2j+1$ of the group SU(2) that implements rotations in quantum mechanics, with basis vectors $|j,m\rangle$, for $m = -j, -j+1, \dots, j$, ordered by m. Put

$$C_{\text{sec}} = \sum_{n\geq 1} \sum_{1\leq i\leq l\leq n} |e_i^n\rangle\langle e_i^n| \otimes |e_l^n\rangle\langle e_l^n|,$$

the projection onto "both systems lie in the same sector, and Alice's level does not exceed Bob's". In the sector $\mathbb{C}^n$, the states $e_1^n, \dots, e_n^n$ form a 1-ladder of length n, so the projection C_{sec} is not uniformly stable and has 1-trees of every finite height. The projection C_{sec} is, however, stable: for a density operator ρ on H_1, the effective observable $T_{C_{\text{sec}}}(\rho)$ that it leaves on Bob's side depends only on the diagonal entries $\rho_{ii}^n = \langle e_i^n|\rho|e_i^n\rangle$, and $T_{C_{\text{sec}}} = J \circ a$ with

$$a\colon \mathcal{T}(H_1) \to \Big(\bigoplus_n \ell^1_n\Big)_{\ell^2}, \quad a(\tau) = \big((\tau^n_{ii})_{i\leq n}\big)_n, \qquad J((w_n)_n) = \bigoplus_n \sum_{l\leq n} \Big(\sum_{i\leq l} w_{n,i}\Big) |e_l^n\rangle\langle e_l^n|,$$

where ℓ^1_n is $\mathbb{C}^n$ with the norm $\|w\|_1 = \sum_i |w_i|$, and $(\bigoplus_n \ell^1_n)_{\ell^2}$ is the space of sequences $(w_n)_n$ with $w_n = (w_{n,i})_{i\leq n}$ in ℓ^1_n and $\sum_n \|w_n\|_1^2 < \infty$, with the norm $(\sum_n \|w_n\|_1^2)^{1/2}$. Here, for $\tau \in \mathcal{T}(H_1)$ with diagonal entries $\tau^n_{ii} = \langle e_i^n|\tau|e_i^n\rangle$, we have $\|a(\tau)\| \leq \sum_n \sum_i |\tau^n_{ii}| \leq \|\tau\|_1$, and $\|J\| \leq 1$, and the ℓ^2-sum of finite-dimensional spaces is reflexive, so Theorem 5.1(4), the factorization of the slice map through a reflexive space, applies.

The space $(\bigoplus_n \ell^1_n)_{\ell^2}$ is reflexive but not superreflexive. Superreflexivity, the reflexivity of all ultrapowers of a space, is relevant here: if the slice map T_C of an observable C factors through a superreflexive space E, then the ultrapower expectation factors through the reflexive space $E^{\mathcal{U}}$, the ultrapower of E along $\mathcal{U}$ [Hei80], so the observable C is uniformly stable. Whether the converse holds is Question (Q2).

9.3 Counting the preparations resolved by adaptive probes

For every positive integer n, a decision tree of height n lets n adaptive probes tell apart its 2^n leaves. We now show that this exponential behavior is the only alternative to polynomial behavior.

Fix $\delta > 0$, and for the moment also a threshold γ and a sign θ. For an integer $n \geq 0$, a *probe tree of depth* n is a family $(\sigma_s)_{s \in 2^{<n}}$ of states of Bob's system. A branch $\eta \in 2^n$ is *realized* by $\rho \in \mathfrak{D}(H_1)$ if, whenever $s^\frown\iota \sqsubseteq \eta$, the inequalities of Definition 9.1(b) hold for $\varphi_C(\rho, \sigma_s)$. Realized branches are the outcomes that the adaptive protocol can return with margin δ, and two preparations that realize different branches are told apart by n probes. For a set $\mathcal{R} \subseteq \mathfrak{D}(H_1)$ of preparations, let $d(\mathcal{R}) \in \{-1, 0, 1, \ldots, \infty\}$ be the supremum of the heights of the δ-trees with these γ and θ whose leaves lie in $\mathcal{R}$, with $d(\emptyset) = -1$. The *growth function* of C is

$$\mathrm{gr}_C(n, \delta) = \sup \#\{\eta \in 2^n : \eta \text{ is realized by some } \rho \in \mathfrak{D}(H_1)\},$$

the supremum being over γ, θ and probe trees of depth n, and the *tree dimension* $d_C(\delta) = \sup_{\gamma,\theta} d(\mathfrak{D}(H_1))$ is the supremum of the heights of the δ-trees for C. We write $\binom{n}{\leq d} = \sum_{i=0}^{d} \binom{n}{i}$.

Theorem 9.9 (Polynomial versus exponential growth). *Let $C \in B(H_1 \otimes H_2)$ be self-adjoint, and let $\delta > 0$ be given.*

(1) *If the tree dimension $d = d_C(\delta)$ is finite, then for every integer $n \geq 0$ we have*

$$\mathrm{gr}_C(n, \delta) \leq \binom{n}{\leq d},$$

a polynomial in n of degree d, and $\binom{n}{\leq d} \leq (\mathrm{e}n/d)^d$ whenever $n \geq d \geq 1$.

(2) *If the tree dimension $d_C(\delta)$ is infinite, then the growth function satisfies $\mathrm{gr}_C(n, \delta) = 2^n$ for every integer $n \geq 0$.*

Consequently, the observable C is uniformly stable if and only if $\mathrm{gr}_C(\cdot, \delta)$ grows polynomially for every $\delta > 0$; otherwise, for some $\delta > 0$, the growth function satisfies $\mathrm{gr}_C(n, \delta) = 2^n$ for every integer $n \geq 0$. In particular, unstable observables have exponential growth.

Proof. Statement (2) holds because a δ-tree of height at least n, restricted to its nodes of length less than n and one leaf above each branch of length n, realizes every branch of a probe tree of depth n. For (1), the idea is that at each probe, at most one of the two halves of the remaining preparations can carry a tree as tall as the whole set. Fix γ and θ. We show by induction on n that, for every probe tree of depth n and every set $\mathcal{R} \subseteq \mathfrak{D}(H_1)$, the number of branches realized by elements of $\mathcal{R}$ is at most $\binom{n}{\leq d(\mathcal{R})}$, with the convention that $\binom{n}{\leq -1} = 0$. For $n = 0$ this is clear. For $n \geq 1$, let $\mathcal{R}_\iota$ be the set of elements of $\mathcal{R}$ that satisfy the inequality for ι at the root. The sets $\mathcal{R}_0$ and $\mathcal{R}_1$ are disjoint, because the two inequalities are separated by δ. A branch $\iota^\frown\eta'$ is realized in $\mathcal{R}$ exactly when η' is realized in $\mathcal{R}_\iota$ for the subtree $(\sigma_{\iota^\frown s})_s$ of depth $n - 1$. We have $d(\mathcal{R}_\iota) \leq d(\mathcal{R})$, and the two numbers $d(\mathcal{R}_0)$ and $d(\mathcal{R}_1)$ cannot both equal $d(\mathcal{R})$: two δ-trees of height $d(\mathcal{R})$ with leaves in $\mathcal{R}_0$ and in $\mathcal{R}_1$, placed above the root probe $\sigma_\emptyset$, would form a δ-tree of height $d(\mathcal{R}) + 1$ with leaves in $\mathcal{R}$. By induction and Pascal's rule, the number of realized branches is at

most $\binom{n-1}{\le d} + \binom{n-1}{\le d-1} = \binom{n}{\le d}$, with $d = d(\mathcal{R})$. The bound $\binom{n}{\le d} \le (en/d)^d$ is standard. The last statements follow from Theorems 9.4 and 9.6. □

Remark 9.10 (Rowbottom's definition of stability). Stability has many equivalent characterizations, but its original definition was a growth condition. Rowbottom [Row64] called a complete theory T *stable in* an infinite cardinal λ if for every model $\mathcal{M}$ of T and every set $A \subseteq \mathcal{M}$ of cardinality λ, the space $S_1(A)$ of complete 1-types over A has cardinality at most λ. The reverse inequality is trivial, since distinct elements of A realize distinct types over A; hence stability in λ says that $|S_1(A)| = |A|$, that is, the number of types grows linearly with the number of parameters, whereas a priori it can be as large as 2^λ. For a single formula $\varphi(x, y)$, stability in T is equivalent to the bound $|S_\varphi(A)| \le |A|$ for every infinite set A of parameters from a model of T, where $S_\varphi(A)$ is the space of complete φ-types over A, the complete descriptions of which instances $\varphi(x, a)$, with a from A, an element satisfies (see [Pil96, Chapter 1] and [She90]). Theorem 9.9 is a finite and adaptive counterpart of this condition: for a uniformly stable observable, the number of preparations resolved by n probes at a fixed margin is bounded by a polynomial in n. For an infinite cardinal λ and an integer $d \ge 1$, we have $\lambda^d = \lambda$, so on infinite parameter sets the distinction between linear and polynomial bounds disappears.

In Theorem 9.11 below, we show that for observables that factor through a Hilbert space, the degree of the polynomial is controlled explicitly, by a linear learner.

Theorem 9.11 (Linear bound for Hilbert-factorable observables). *Let $C \in B(H_1 \otimes H_2)$ be self-adjoint, and suppose that there are a Hilbert space $\mathcal{H}$, positive real numbers R_u and R_v, and maps $u\colon \mathfrak{D}(H_1) \to \mathcal{H}$ and $v\colon \mathfrak{D}(H_2) \to \mathcal{H}$ such that $\|u(\rho)\| \le R_u$, $\|v(\sigma)\| \le R_v$ and $\varphi_C(\rho, \sigma) = \langle u(\rho), v(\sigma)\rangle$ for all $\rho \in \mathfrak{D}(H_1)$ and $\sigma \in \mathfrak{D}(H_2)$. Then for every $\delta > 0$, we have*

$$d_C(\delta) \le 16\Big(\frac{R_u R_v}{\delta}\Big)^2,$$

independently of the dimensions of H_1 and H_2. In particular, the tree dimension is at most $16/\delta^2$ for the swap operator F and for the equality tests $[A = B]$ of Section 4, and at most $4/\delta^2$ for the swap test P_+ and the Hong–Ou–Mandel projection P_-.

Proof. The idea is to run the perceptron, the algorithm that updates a linear classifier by adding to it each misclassified example multiplied by the label ± 1 of that example, along the tree, always moving to the child that contradicts its prediction; the leaf reached at the end is a linear separator with margin $\delta/2$, and the perceptron makes few mistakes against such a separator. Regard the Hilbert space as a real inner product space with $\langle\cdot,\cdot\rangle_{\mathbb{R}} = \mathrm{Re}\langle\cdot,\cdot\rangle$; since φ_C is real, we have $\varphi_C = \langle u, v\rangle_{\mathbb{R}}$. Let a δ-tree of height n be given, with threshold γ and sign θ. We may assume that $n \ge 1$; the tree then requires values of φ_C on both sides of γ, which forces $|\gamma| \le R_u R_v$. Put $\hat{u}(\rho) = (u(\rho), R_u)$ and $\hat{v}(\sigma) = (v(\sigma), -\gamma/R_u)$ in $\mathcal{H} \oplus \mathbb{R}$; then $\langle \hat{u}(\rho), \hat{v}(\sigma)\rangle_{\mathbb{R}} = \varphi_C(\rho, \sigma) - \gamma$, $\|\hat{u}\| \le \sqrt{2}R_u$ and $\|\hat{v}\| \le \sqrt{2}R_v$. Start with $z_0 = 0$ at the root $s_0 = \emptyset$. At the node s_i reached at step i, for $0 \le i < n$, let $\hat{y}_i = \mathrm{sign}\langle z_i, \hat{v}(\sigma_{s_i})\rangle_{\mathbb{R}}$, with $\mathrm{sign}\, 0 = 1$, move to the child $s_{i+1} = s_i{}^\frown\iota$ whose label $y_i = \theta(2\iota - 1)$ equals $-\hat{y}_i$, and put $z_{i+1} = z_i + y_i\,\hat{v}(\sigma_{s_i})$. After n steps we reach a branch η, and its leaf $\hat{u}^* = \hat{u}(\rho_\eta)$ satisfies $y_i\langle \hat{u}^*, \hat{v}(\sigma_{s_i})\rangle_{\mathbb{R}} \ge \delta/2$ for every $i < n$. Summing over the steps, we obtain $\langle \hat{u}^*, z_n\rangle_{\mathbb{R}} \ge n\delta/2$, while $\|z_n\|^2 \le \sum_i \|\hat{v}(\sigma_{s_i})\|^2 \le 2nR_v^2$, because each step adds $2y_i\langle z_i, \hat{v}(\sigma_{s_i})\rangle_{\mathbb{R}} \le 0$ to the cross term. By the Cauchy–Schwarz inequality, we have $n\delta/2 \le \sqrt{2}R_u \cdot \sqrt{2n}\,R_v$, that is, $n \le 16(R_u R_v/\delta)^2$.

For F and $[A = B]$, both radii R_u and R_v are at most 1 (see Section 4.3 and Proposition 4.6). Since $\varphi_{P_\pm} = \frac{1}{2}(1 \pm \varphi_F)$, every δ-tree for $P_\pm$ is a 2δ-tree for F. □

Remark 9.12 (Node-dependent thresholds). The proof of Theorem 9.11 uses the threshold only through the probe vectors $\hat{v}(\sigma_s)$, and the bound $|\gamma| \le R_u R_v$ holds at each node separately, since every node has leaves on both sides of its threshold. The proof therefore applies without change to trees in which the threshold γ_s depends on the node s. Consequently, for a Hilbert-factorable observable, the response class $\{\sigma \mapsto \varphi_C(\rho, \sigma) : \rho \in \mathfrak{D}(H_1)\}$ has sequential fat-shattering dimension at most $16(R_u R_v/\delta)^2$ at scale δ in the sense of Rakhlin, Sridharan and Tewari [RST15], that is, δ-trees for C with node-dependent thresholds have height at most this number, and the class is online learnable in that sense, with bounds independent of the dimension.

9.4 Affine approximation versus exponential growth

The following table collects the statements proved in this paper that correspond, for bipartite observables, to the dichotomy of deep computations described in the Introduction (see [ADIW24, DIMW+26]).

	stable side	unstable side
idealized responses	limits of *affine* combinations of actual responses, that is, of finite mixtures (see Theorem 6.1)	for the order projection, not approximable within any error less than 1 (see Example 6.3)
preparations resolved by n adaptive probes at margin δ	$O(n^{d_C(\delta)})$ if C is uniformly stable (see Theorem 9.9)	2^n at some margin (see Theorem 9.4)
Hilbert-factorable observables	for factors u and v on states bounded by 1: tree dimension at most $16/\delta^2$, and $4/\delta^2$ for the swap test and Hong–Ou–Mandel coincidences; for a linear and a conjugate-linear factor of the bilinear form, bounded by 1: mixtures of $\lceil 4/\varepsilon^2 \rceil$ preparations simulate idealized preparations within ε, for every $\varepsilon > 0$ (see Theorems 9.11 and 6.5)	—

The approximating combinations in the first row are affine because expectations are affine in the state (see Remark 6.4); the polynomial in the second row has degree $d_C(\delta)$, which need not be one: for the comparison GT_r of Section 11, the projection onto the event that, on two registers of r qubits read as binary numbers, Alice's number exceeds Bob's, the degree is at least $r - 1$ at margin 1. In the Hilbert-factorable case, linearity enters twice more: idealized responses are inner products with a fixed vector, and tree heights are bounded by a linear classifier. The first two rows concern different notions. The sector comparison of Example 9.8 is stable and therefore admits affine simulation of idealized preparations; it is not uniformly stable, however, and its growth at margin 1 is 2^n.

9.5 Learning: Littlestone's mistake bound

The δ-trees of Definition 9.1 also govern mistake bounds in online learning [Lit88]. Fix a threshold γ and a margin $\delta > 0$, and consider the following game. Alice prepares an unknown state ρ. In each round, Bob chooses a probe σ for which $|\varphi_C(\rho, \sigma) - \gamma| \geq \delta/2$; a learner predicts whether $\varphi_C(\rho, \sigma)$ lies above or below γ, and is then told the answer. Let $\mathrm{mist}_C(\gamma, \delta)$ be the least integer m such that some deterministic learner makes at most m mistakes, for every $\rho \in \mathfrak{D}(H_1)$ and every sequence of probes, or ∞ if there is no such integer, and let $\mathrm{mist}_C(\delta) = \sup_\gamma \mathrm{mist}_C(\gamma, \delta)$. Let d_γ be the supremum of the heights of the δ-trees for C with threshold γ and sign $\theta = 1$. Since reflecting a tree, that is, exchanging the two children at every node, exchanges $\theta = 1$ and $\theta = -1$, we have $\sup_\gamma d_\gamma = d_C(\delta)$.

Proposition 9.13 (Mistake bound equals tree dimension). *For every self-adjoint $C \in B(H_1 \otimes H_2)$, every real number γ and every $\delta > 0$, the mistake bound $\mathrm{mist}_C(\gamma, \delta)$ equals d_γ, the supremum of the heights of the δ-trees for C with threshold γ and sign $\theta = 1$. Consequently, the observable C is uniformly stable if and only if $\mathrm{mist}_C(\delta)$ is finite for every $\delta > 0$, and then $\mathrm{mist}_C(\delta) = d_C(\delta)$.*

Proof. If there is a tree of height n, an adversary walks down it, choosing at each node the child that contradicts the learner's prediction; the leaf reached is a preparation consistent with all answers, so every learner makes n mistakes. Conversely, run Littlestone's standard optimal algorithm: keep the set V of preparations consistent with the answers so far; for a probe σ, let V^1 and V^0 be the sets of elements of V above $\gamma + \delta/2$ and below $\gamma - \delta/2$, and predict the side ι for which $d(V^\iota)$ is larger, in the notation of Section 9.3. As in the proof of Theorem 9.9, the numbers $d(V^0)$ and $d(V^1)$ cannot both equal $d(V)$, so after a mistake the new consistent set $V^{1-\iota}$ satisfies $d(V^{1-\iota}) < d(V)$. Since the true preparation stays in V, we have $d(V) \geq 0$ throughout, and there are at most d_γ mistakes. The second statement follows from Theorem 9.6. □

By Theorem 9.4 and Proposition 9.13, for an unstable observable the mistake bound in the threshold game is infinite at some margin, and the sector comparison of Example 9.8 shows that stability alone does not guarantee a finite mistake bound. The threshold game is a restricted form of online learning: general online learnability of real-valued classes allows thresholds that vary from probe to probe (see the historical remarks below and Question (Q6)). In the Hilbert-factorable case the restriction is immaterial: by Remark 9.12, the perceptron of Theorem 9.11 learns online, even when the thresholds vary and in every dimension, with at most $16(R_u R_v/\delta)^2$ mistakes, where R_u and R_v bound the norms of the two factors u and v of the expectation.

9.6 Historical remarks

Stability in a cardinal λ, which extends Morley's ω-stability (stability in $\aleph_0$) [Mor65] to uncountable languages, was introduced by Rowbottom [Row64] in his work on the Łoś conjecture for uncountable theories, the conjecture that a complete theory with exactly one model, up to isomorphism, in some cardinal larger than the cardinality of its language has exactly one in every such cardinal. Shelah studied the stability spectrum, the class of cardinals in which a theory is stable, and related it to the order property [She71, She90]. For Banach space structures, the stability spectrum theorem was proved by Iovino in his dissertation [Iov94]. Binary trees of formulas and the associated rank are due to Shelah [She71, She90], and the equivalence between bounded ladders, bounded trees and stability of a formula in all models of a theory, which underlies Theorem 9.6, is part

of classical stability theory [She90]. For Banach spaces, the stability of all ultrapowers is Raynaud's superstability [Ray83], the analytic counterpart of condition (3) of Theorem 9.6; Proposition 9.3 adapts Proposition 2.3.1 of Iovino's notes on stability theory [Iov], where it is proved for first-order formulas, with the gap δ given by the failure of the double limit property in place of the negation of a formula; in Theorem 9.6, ultrapowers take the place of compactness. Quantitative bounds in the converse direction, from trees back to ladders, come from a Ramsey argument for trees due to Hodges [Hod93]; Conant and Terry give explicit bounds for real-valued functions [CT26]. Theorem 9.9 is the tree analogue of the Sauer–Shelah lemma [Sau72, She72]; versions for Littlestone classes appear in the work of Ben-David, Pál and Shalev-Shwartz [BDPSS09] and of Rakhlin, Sridharan and Tewari [RST15]. Littlestone introduced the dimension and the optimal algorithm that bear his name [Lit88]; for real-valued classes, online learnability is characterized by finiteness of the sequential fat-shattering dimension at every scale [RST15], which allows node-dependent thresholds (see Question (Q6)). The link between stability and Littlestone dimension in model theory is due to Chase and Freitag [CF19]. The bound in Theorem 9.11 is Novikoff's mistake bound for the perceptron [Nov62, Blo62], read as a bound on tree height. For n-qubit states probed by two-outcome measurements, Aaronson, Chen, Hazan, Kale and Nayak showed that the sequential fat-shattering dimension is $O(n)$ at each fixed scale [ACH+18].

10 Order: stable observables are closed under norm limits and orthogonal sums

We now determine the algebraic and topological operations under which stability is preserved, and locate the boundary of these closure properties: we show that stability survives norm limits and orthogonal sums, but neither meets, which are the projections onto the intersections of two ranges, nor limits in the strong operator topology, which are the pointwise limits on vectors.

10.1 Topological and algebraic closure

Proposition 10.1 (Topological and algebraic closure). *The set $\mathcal{O}$ of stable observables is a norm-closed operator system: a self-adjoint linear subspace of $B(H_1 \otimes H_2)$ that contains I and is closed in the norm topology. The set $\mathcal{O}$ contains the minimal tensor product $B(H_1) \otimes_{\min} B(H_2)$ and every observable whose expectation factors through a Hilbert space. The same holds for the set of uniformly stable observables.*

Proof. The identity $I = I \otimes I$ is stable, and for density operators we have $\varphi_{C^*} = \overline{\varphi_C}$. Linear combinations and norm limits are handled by Lemma 3.4(a) and (b), and for uniform stability the same lemma applies to the ultrapower expectation $\Phi_C^{\mathcal{U}}$ of Section 9.2. The last two statements are Theorem 4.1, Proposition 4.4 and Proposition 9.7. □

Ordered by inclusion of ranges, with $I-P$ as the orthocomplement of P, the projections on $H_1 \otimes H_2$ form an *orthomodular poset*: any two mutually orthogonal projections P and Q, that is, projections with $PQ = 0$, have a join, namely $P + Q$, and whenever $P \leq Q$, the projection Q is the join of P and a projection orthogonal to P, namely $Q - P$; this last property is the *orthomodular law*.

Corollary 10.2 (Orthocomplements and orthogonal sums). *The stable projections form an orthomodular poset: if a projection P on $H_1 \otimes H_2$ is stable, then so is $I - P$, and*

if two stable projections P and Q on $H_1 \otimes H_2$ satisfy $PQ = 0$, then the projection $P + Q$ is stable.

Proof. If the projections P and Q are stable, then the projections $I - P$, $P + Q$ (when $PQ = 0$) and $Q - P$ (when $P \leq Q$) are linear combinations of stable observables, and hence they are stable by Proposition 10.1. Thus the stable projections are closed under orthocomplements and orthogonal sums, and the orthomodular law, valid for all projections on $H_1 \otimes H_2$, holds among them. □

10.2 The boundary of closure: meets and strong limits

Closure first fails at meets: in Proposition 10.3 below, we construct two uniformly stable projections whose meet is unstable.

Proposition 10.3 (The boundary: meets). *If H_1 and H_2 are infinite-dimensional, then there are projections P_1 and P_2 in $B(H_1) \otimes_{\min} B(H_2)$, hence uniformly stable, whose meet $P_1 \wedge P_2$, the projection onto the intersection of their ranges, is unstable.*

Proof. The idea is to hide the order projection inside the intersection of two subspaces that are compact perturbations of a product projection, by tilting each matrix entry into a fresh direction through a small angle, in opposite senses for the two subspaces. Let $(e_k)_{k\geq 1}$ and $(e'_k)_{k\geq 1}$ be sequences in H_1 whose terms together form an orthonormal set, let $(f_l)_{l\geq 1}$ and $(f'_l)_{l\geq 1}$ be such sequences in H_2, and let $P_{k<l}$ be the order projection built from (e_k) and (f_l), the projection onto the closed span of the vectors $e_k \otimes f_l$ with $k < l$. For $k \geq l$, let $\zeta_k = 1/k$ and

$$n_{kl} = \cos\zeta_k\, e_k \otimes f_l + \sin\zeta_k\, e'_k \otimes f'_l, \qquad n'_{kl} = \cos\zeta_k\, e_k \otimes f_l - \sin\zeta_k\, e'_k \otimes f'_l,$$

and let P_1 and P_2 be the projections onto the closed subspaces $\operatorname{ran} P_{k<l} \oplus \overline{\operatorname{span}}\{n_{kl} : k \geq l\}$ and $\operatorname{ran} P_{k<l} \oplus \overline{\operatorname{span}}\{n'_{kl} : k \geq l\}$. The planes spanned by $e_k \otimes f_l$ and $e'_k \otimes f'_l$ are mutually orthogonal and orthogonal to the range of $P_{k<l}$, and in each of them the lines through n_{kl} and n'_{kl} meet only in 0, since $0 < \zeta_k < \pi/2$. Both subspaces are block diagonal for this decomposition, so their intersection is the range of $P_{k<l}$, that is, $P_1 \wedge P_2 = P_{k<l}$, which is unstable by Theorem 4.7.

Let Q_1 and Q_2 be the projections onto the closed spans of (e_k) and (f_l). The operator $P_1 - Q_1 \otimes Q_2$ is the orthogonal sum over $k \geq l$ of the rank-two operators $|n_{kl}\rangle\langle n_{kl}| - |e_k \otimes f_l\rangle\langle e_k \otimes f_l|$, of norm $\sin\zeta_k$. For every $\varepsilon > 0$, only finitely many pairs $k \geq l$ have $1/k \geq \varepsilon$, so this sum is compact. Since $Q_1 \otimes Q_2$ and the compact operators lie in $B(H_1) \otimes_{\min} B(H_2)$, so does P_1, and likewise P_2. Conclude with Theorem 4.1 and Proposition 9.7. □

By von Neumann's alternating projection theorem, the meet is the strong limit $P_1 \wedge P_2 = \lim_{m\to\infty} (P_1 P_2)^m$ [Deu01]. The convergence is in norm if and only if the sum of the ranges of P_1 and P_2 is closed [Deu01]. In Proposition 10.3, the angles ζ_k accumulate at 0, so the sum of the ranges is not closed. Since $B(H_1) \otimes_{\min} B(H_2)$ is an algebra, the operators $(P_1 P_2)^m$, for positive integers m, are stable, and their strong limit $P_1 \wedge P_2$ is unstable. Thus stability passes to norm limits (see Lemma 3.4(b)), but not to strong limits.

11 Outlook: registers of growing size

We indicate, finally, how the uniform notions of Section 9 apply to quantum computation. There the natural setting is a family of registers of growing size, each with a readout

observable on two copies of the register, and deep computations are idealized as ultralimits of the outputs of longer and longer compositions of quantum channels, the linear maps that send density operators to density operators even when they act on one part of a larger system, in the sense of the work on deep computations [ADIW24, DIMW+26].

To make this precise, for every positive integer n, let $H^{(n)} = (\mathbb{C}^2)^{\otimes n}$ be the register of n qubits, or two-level systems, and let C_n be a readout observable on $H^{(n)} \otimes H^{(n)}$; the outputs of deep computations are the ultralimit states $\lim_{k,\mathcal{U}} \mathcal{E}_k \circ \cdots \circ \mathcal{E}_1(\rho_0)$, that is, the states $A \mapsto \lim_{k,\mathcal{U}} \operatorname{Tr}(\mathcal{E}_k \circ \cdots \circ \mathcal{E}_1(\rho_0)\, A)$, along a nonprincipal ultrafilter $\mathcal{U}$ on $\mathbb{N}$, for an initial density operator ρ_0 on a register and a sequence $(\mathcal{E}_k)_{k\in\mathbb{N}}$ of quantum channels on that register. On a single register, which is finite-dimensional, such a state is again given by a density operator. Each C_n acts on a finite-dimensional space and is thus stable (see Corollary 4.2); the meaningful condition is the uniform stability of the family, in the sense of the paragraph after Theorem 9.6.

Two families illustrate the distinction. The equality tests $\mathrm{EQ}_n = \sum_x |x\rangle\langle x| \otimes |x\rangle\langle x|$, where $x = x_1 \cdots x_n$ runs over the strings in $\{0,1\}^n$ and $|x\rangle = |x_1\rangle \otimes \cdots \otimes |x_n\rangle$ is the vector that x labels in the computational basis of $H^{(n)}$, the product basis built from a fixed orthonormal basis $|0\rangle, |1\rangle$ of $\mathbb{C}^2$, factor through ℓ^2 with constants independent of n (see the proof of Proposition 4.6). The family is therefore uniformly stable, its tree dimension at every margin $\delta > 0$ is at most $16/\delta^2$ for all n (see Theorem 9.11), so its growth is polynomial of a degree independent of n, and the Monte Carlo bound of Corollary 6.6 holds, at every accuracy $\varepsilon > 0$, with mixtures of $M = \lceil 4/\varepsilon^2 \rceil$ preparations for all n. The comparisons $\mathrm{GT}_n = \sum_{x>y} |x\rangle\langle x| \otimes |y\rangle\langle y|$, in which the strings x and y are compared as binary numbers, contain 1-ladders of length $2^n - 1$, hence 1-trees of height $n - 1$ by Proposition 9.3: the family is not uniformly stable, and $n - 1$ adaptive probes resolve 2^{n-1} basis preparations. The same dichotomy appears in communication complexity, the study of how many bits two parties holding separate inputs must exchange to compute a function of both: equality has constant public-coin randomized cost, that is, parties who share random bits and may err with small probability can compute it by exchanging a number of bits independent of n, while Greater-Than on n-bit inputs has cost $\Theta(\log n)$ [KN97, Vio15]. In the language of graphs, the communication matrix of Greater-Than, the table of its values on all pairs of inputs, is the biadjacency matrix of a *half-graph*: its entries form a staircase of 0s and 1s, like the values of a ladder in the sense of Definition 9.1. A family of communication problems excludes large Greater-Than instances exactly when its matrices contain no large half-graphs. Harms, Wild and Zamaraev showed that this exclusion, which is stability, is necessary for constant randomized cost [HWZ22], and Harms and Zamaraev proved it sufficient for matrices of sign-rank at most 3, those that record the signs of the entries of a real matrix of rank at most 3, and for unit disk graphs, the intersection graphs of disks of equal radius in the plane [HZ23].

12 Questions

We close with questions raised by the results above, on the structure of the operator system of stable observables, on the gap between stability and uniform stability, and on quantitative bounds:

(Q1) Is every stable observable a norm limit of observables whose slice maps, which send each preparation of Alice's system to the effective observable that it leaves on Bob's side, factor through a Hilbert space?

(Q2) Is a self-adjoint observable C uniformly stable if and only if its slice map T_C factors through a superreflexive Banach space, one all of whose ultrapowers are reflexive?

(Q3) Is the operator system $\mathcal{O}$ of stable observables closed under products, and hence a C*-algebra? If so, the meet of two stable projections whose ranges have a closed sum would be stable.

(Q4) Is the meet of two *commuting* stable projections stable?

(Q5) For a uniformly stable family of readouts on n-qubit registers, what is the minimal size of a mixture simulating idealized preparations to precision ε, and when is it independent of n? For Hilbert-factorable families, Theorem 6.5 gives a bound independent of n.

(Q6) Theorem 9.9 and Proposition 9.13 concern trees with a constant threshold. Does uniform stability of a self-adjoint observable C imply that its response class $\{\sigma \mapsto \varphi_C(\rho, \sigma) : \rho \in \mathfrak{D}(H_1)\}$ has finite sequential fat-shattering dimension at every scale, with node-dependent thresholds? Remark 9.12 gives a positive answer for Hilbert-factorable observables.

(Q7) Is the tree dimension $d_C(\delta)$ of a uniformly stable self-adjoint observable C, the supremum of the heights of its δ-trees, always bounded by a polynomial in $1/\delta$, as in the Hilbert case, and is the degree of that polynomial related to the Rademacher type of a space through which the slice map T_C factors?

References

[ACH+18] S. Aaronson, X. Chen, E. Hazan, S. Kale, and A. Nayak, *Online learning of quantum states*, Advances in Neural Information Processing Systems 31 (NeurIPS 2018), 2018, pp. 8962–8972.

[ADIW24] Samson Alva, E. Dueñez, J. N. Iovino, and C. Walton, *Approximability of deep computations*, 2024, Accepted; to appear in Mathematical Structures in Computer Science. arXiv:2409.06064.

[Are51] R. Arens, *The adjoint of a bilinear operation*, Proc. Amer. Math. Soc. **2** (1951), 839–848.

[BCWdW01] H. Buhrman, R. Cleve, J. Watrous, and R. de Wolf, *Quantum fingerprinting*, Phys. Rev. Lett. **87** (2001), 167902.

[BDPSS09] S. Ben-David, D. Pál, and S. Shalev-Shwartz, *Agnostic online learning*, Proc. 22nd Conference on Learning Theory (COLT 2009), 2009.

[Bir35] G. Birkhoff, *Abstract linear dependence and lattices*, Amer. J. Math. **57** (1935), 800–804.

[BL71] J. T. Baldwin and A. H. Lachlan, *On strongly minimal sets*, J. Symbolic Logic **36** (1971), 79–96.

[Blo62] H. D. Block, *The perceptron: a model for brain functioning. I*, Rev. Modern Phys. **34** (1962), 123–135.

[BY14] Itaï Ben Yaacov, *Model theoretic stability and definability of types, after A. Grothendieck*, Bull. Symb. Log. **20** (2014), no. 4, 491–496. MR 3294276

[BYBHU08] Itaï Ben Yaacov, Alexander Berenstein, C. Ward Henson, and Alexander Usvyatsov, *Model theory for metric structures*, Model theory with applications to algebra and analysis. Vol. 2, London Math. Soc. Lecture Note Ser., vol. 350, Cambridge Univ. Press, Cambridge, 2008, pp. 315–427. MR MR2436146

[BYU10] Itaï Ben Yaacov and Alexander Usvyatsov, *Continuous first order logic and local stability*, Trans. Amer. Math. Soc. **362** (2010), no. 10, 5213–5259. MR 2657678

[CF19] H. Chase and J. Freitag, *Model theory and machine learning*, Bull. Symbolic Logic **25** (2019), 319–332.

[CR70] H. H. Crapo and G.-C. Rota, *On the Foundations of Combinatorial Theory: Combinatorial Geometries*, preliminary ed., MIT Press, Cambridge, MA, 1970.

[CT26] G. Conant and C. Terry, *Encoding orders and trees in real-valued functions*, arXiv:2607.21761, 2026.

[Deu01] F. Deutsch, *Best Approximation in Inner Product Spaces*, CMS Books in Mathematics, Springer, New York, 2001.

[DFJP74] W. J. Davis, T. Figiel, W. B. Johnson, and A. Pełczyński, *Factoring weakly compact operators*, J. Functional Analysis **17** (1974), 311–327.

[DIMW+26] E. Dueñez, J. Iovino, T. Matos-Wiederhold, L. Salvetti, and F. D. Tall, *Complexity of deep computations via topology of function spaces*, 2026, In preparation. arXiv:2601.00528.

[Ebe47] W. F. Eberlein, *Weak compactness in Banach spaces. I*, Proc. Nat. Acad. Sci. U.S.A. **33** (1947), 51–53.

[Gan40] V. Gantmacher, *Über schwache totalstetige Operatoren*, Mat. Sbornik N.S. **7(49)** (1940), 301–308.

[Gro52] A. Grothendieck, *Critères de compacité dans les espaces fonctionnels généraux*, Amer. J. Math. **74** (1952), 168–186. MR 47313

[Hei80] S. Heinrich, *Ultraproducts in Banach space theory*, J. Reine Angew. Math. **313** (1980), 72–104.

[HH84] V. Harnik. and L. Harrington, *Fundamentals of forking*, Ann. Pure Appl. Logic **26** (1984), 245–286.

[Hod93] Wilfrid Hodges, *Model theory*, Encyclopedia of Mathematics and its Applications, vol. 42, Cambridge University Press, Cambridge, 1993. MR 94e:03002

[HOM87] C. K. Hong, Z. Y. Ou, and L. Mandel, *Measurement of subpicosecond time intervals between two photons by interference*, Phys. Rev. Lett. **59** (1987), 2044–2046.

[HP11] E. Hrushovski and A. Pillay, *On NIP and invariant measures*, J. Eur. Math. Soc. **13** (2011), 1005–1061.

[HPP08] Ehud Hrushovski, Ya'acov Peterzil, and Anand Pillay, *Groups, measures, and the NIP*, J. Amer. Math. Soc. **21** (2008), no. 2, 563–596. MR 2373360 (2008k:03078)

[HWZ22] N. Harms, S. Wild, and V. Zamaraev, *Randomized communication and implicit graph representations*, STOC '22: Proc. 54th Annual ACM Symposium on Theory of Computing, 2022.

[HZ23] N. Harms and V. Zamaraev, *Randomized communication and implicit representations for matrices and graphs of small sign-rank*, arXiv:2307.04441, 2023.

[Iov] J. Iovino, *Notes on stability theory*, preliminary notes, The University of Texas at San Antonio.

[Iov94] José N. Iovino, *Stable theories in functional analysis*, Ph.D. thesis, University of Illinois at Urbana-Champaign, 1994.

[Iov99a] ———, *Stable Banach spaces and Banach space structures. I. Fundamentals*, Models, algebras, and proofs (Bogotá, 1995), Dekker, New York, 1999, pp. 77–95. MR 2000h:03064a

[Iov99b] ———, *Stable models and reflexive Banach spaces*, J. Symbolic Logic **64** (1999), no. 4, 1595–1600. MR 2001e:03068

[Kei87] H. J. Keisler, *Measures and forking*, Ann. Pure Appl. Logic **34** (1987), 119–169.

[Kei23] H. Jerome Keisler, *Model theory for real-valued structures*, Beyond First-Order Model Theory, Vol. 2 (José Iovino, ed.), CRC Press, 2023, pp. 43–101.

[KM81] Jean-Louis Krivine and Bernard Maurey, *Espaces de Banach stables*, Israel J. Math. **39** (1981), no. 4, 273–295. MR 83a:46030

[KN97] E. Kushilevitz and N. Nisan, *Communication Complexity*, Cambridge University Press, 1997.

[KP97] Byunghan Kim and Anand Pillay, *Simple theories*, Ann. Pure Appl. Logic **88** (1997), no. 2-3, 149–164, Joint AILA-KGS Model Theory Meeting (Florence, 1995). MR MR1600895 (99b:03049)

[KR83] R. V. Kadison and J. R. Ringrose, *Fundamentals of the Theory of Operator Algebras, Vol. I*, Academic Press, 1983.

[KŠ40] M. Krein and V. Šmulian, *On regularly convex sets in the space conjugate to a Banach space*, Ann. of Math. **41** (1940), 556–583.

[Lev07] E. Levy, *Weakly compact "matrices", Fubini-like property and extension of densely defined semigroups of operators*, arXiv:0704.3558, 2007.

[Lit88] N. Littlestone, *Learning quickly when irrelevant attributes abound: a new linear-threshold algorithm*, Machine Learning **2** (1988), 285–318.

[LP79] D. Lascar and B. Poizat, *An introduction to forking*, J. Symbolic Logic **44** (1979), 330–350.

[Maz33] S. Mazur, *Über konvexe Mengen in linearen normierten Räumen*, Studia Math. **4** (1933), 70–84.

[ML36] S. Mac Lane, *Some interpretations of abstract linear dependence in terms of projective geometry*, Amer. J. Math. **58** (1936), 236–240.

[Mor65] M. Morley, *Categoricity in power*, Trans. Amer. Math. Soc. **114** (1965), 514–538.

[NK09] H. Nishimura and S. Kuroda (eds.), *A Lost Mathematician, Takeo Nakasawa: The Forgotten Father of Matroid Theory*, Birkhäuser, Basel, 2009.

[Nov62] A. B. J. Novikoff, *On convergence proofs on perceptrons*, Proc. Symposium on the Mathematical Theory of Automata, vol. 12, Polytechnic Institute of Brooklyn, 1962, pp. 615–622.

[Pil96] A. Pillay, *Geometric stability theory*, The Clarendon Press Oxford University Press, New York, 1996, Oxford Science Publications.

[Pis81] G. Pisier, *Remarques sur un résultat non publié de B. Maurey*, Séminaire d'Analyse Fonctionnelle 1980–1981, Exp. V, École Polytechnique, Palaiseau, 1981.

[Ptá63] V. Pták, *A combinatorial lemma on the existence of convex means and its application to weak compactness*, Convexity, Proc. Sympos. Pure Math., vol. 7, Amer. Math. Soc., 1963, pp. 437–450.

[Pym65] J. S. Pym, *The convolution of functionals on spaces of bounded functions*, Proc. London Math. Soc. (3) **15** (1965), 84–104.

[Ray83] Y. Raynaud, *Espaces de Banach superstables, distances stables et homéomorphismes uniformes*, Israel J. Math. **44** (1983), no. 1, 33–52. MR 84j:46024

[Row64] F. Rowbottom, *The łoś conjecture for uncountable theories*, Notices of the Amer. Math. Soc. **11** (1964), 248–249.

[RST15] A. Rakhlin, K. Sridharan, and A. Tewari, *Online learning via sequential complexities*, J. Mach. Learn. Res. **16** (2015), 155–186.

[Sau72] N. Sauer, *On the density of families of sets*, J. Combinatorial Theory Ser. A **13** (1972), 145–147. MR 307902

[She71] Saharon Shelah, *Stability, the f.c.p., and superstability; model theoretic properties of formulas in first order theory*, Ann. Math. Logic **3** (1971), no. 3, 271–362. MR 0317926

[She72] ———, *A combinatorial problem; stability and order for models and theories in infinitary languages*, Pacific J. Math. **41** (1972), 247–261. MR 307903

[She90] S. Shelah, *Classification theory and the number of nonisomorphic models*, second ed., Studies in Logic and the Foundations of Mathematics, vol. 92, North-Holland Publishing Co., Amsterdam, 1990. MR MR1083551 (91k:03085)

[Sim68] S. Simons, *The iterated limit condition, a Fubini theorem and weak compactness*, Math. Ann. **176** (1968), 87–95.

[Ste10] E. Steinitz, *Algebraische Theorie der Körper*, J. Reine Angew. Math. **137** (1910), 167–309.

[Ü87] A. Ülger, *Weakly compact bilinear forms and Arens regularity*, Proc. Amer. Math. Soc. **101** (1987), 697–704.

[vdW30] B. L. van der Waerden, *Moderne Algebra*, Springer, Berlin, 1930.

[Vio15] E. Viola, *The communication complexity of addition*, Combinatorica **35** (2015), 703–747.

[Š40] V. Šmulian, *Über lineare topologische Räume*, Mat. Sbornik N.S. **7(49)** (1940), 425–448.

[Whi35] H. Whitney, *On the abstract properties of linear dependence*, Amer. J. Math. **57** (1935), 509–533.